

# Collaborating in Motion

Distributed and Stochastic Algorithms for Emergent Behavior  
in Programmable Matter

**Joshua J. Daymude**

Prof. Andréa W. Richa (Chair)

Prof. Christian Scheideler, Prof. Dana Randall, Prof. Theodore Pavlic, and Prof. Stephanie Gil

PhD Dissertation Defense — March 15, 2021

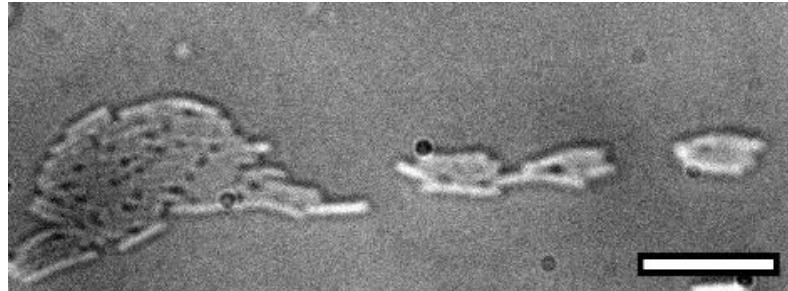
CIDSE, Arizona State University



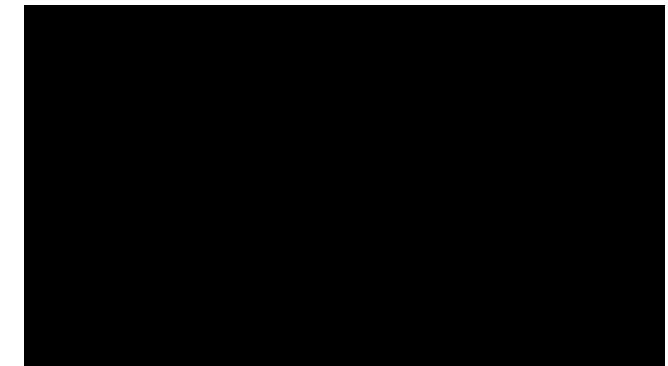
# Self-Organizing Systems

---

Cooperative decentralized systems are capable of surprising emergent behavior arising from relatively simple interactions of their members.



[HMSKCLCA 2011](#)



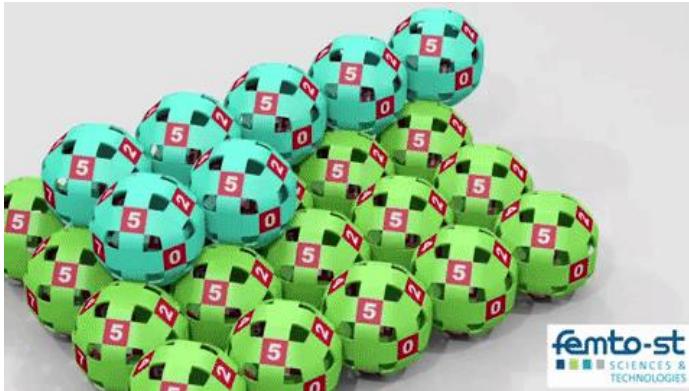
[Microsoft Research 2016](#)

# Programmable Matter

---

**Programmable matter** is a substance that can change its physical properties **autonomously** based on **user input** or **environmental stimuli**.

"Catoms"  
[PB 2018](#)



"Kilobots"  
[RCN 2014](#)



"M-Blocks"  
[RGR 2013](#)



"Particle Robots"  
[LBBCRHRL 2019](#)



# Programmable Matter

---

Centimeter/millimeter-scale robots are more limited than, say, Spot from Boston Dynamics.

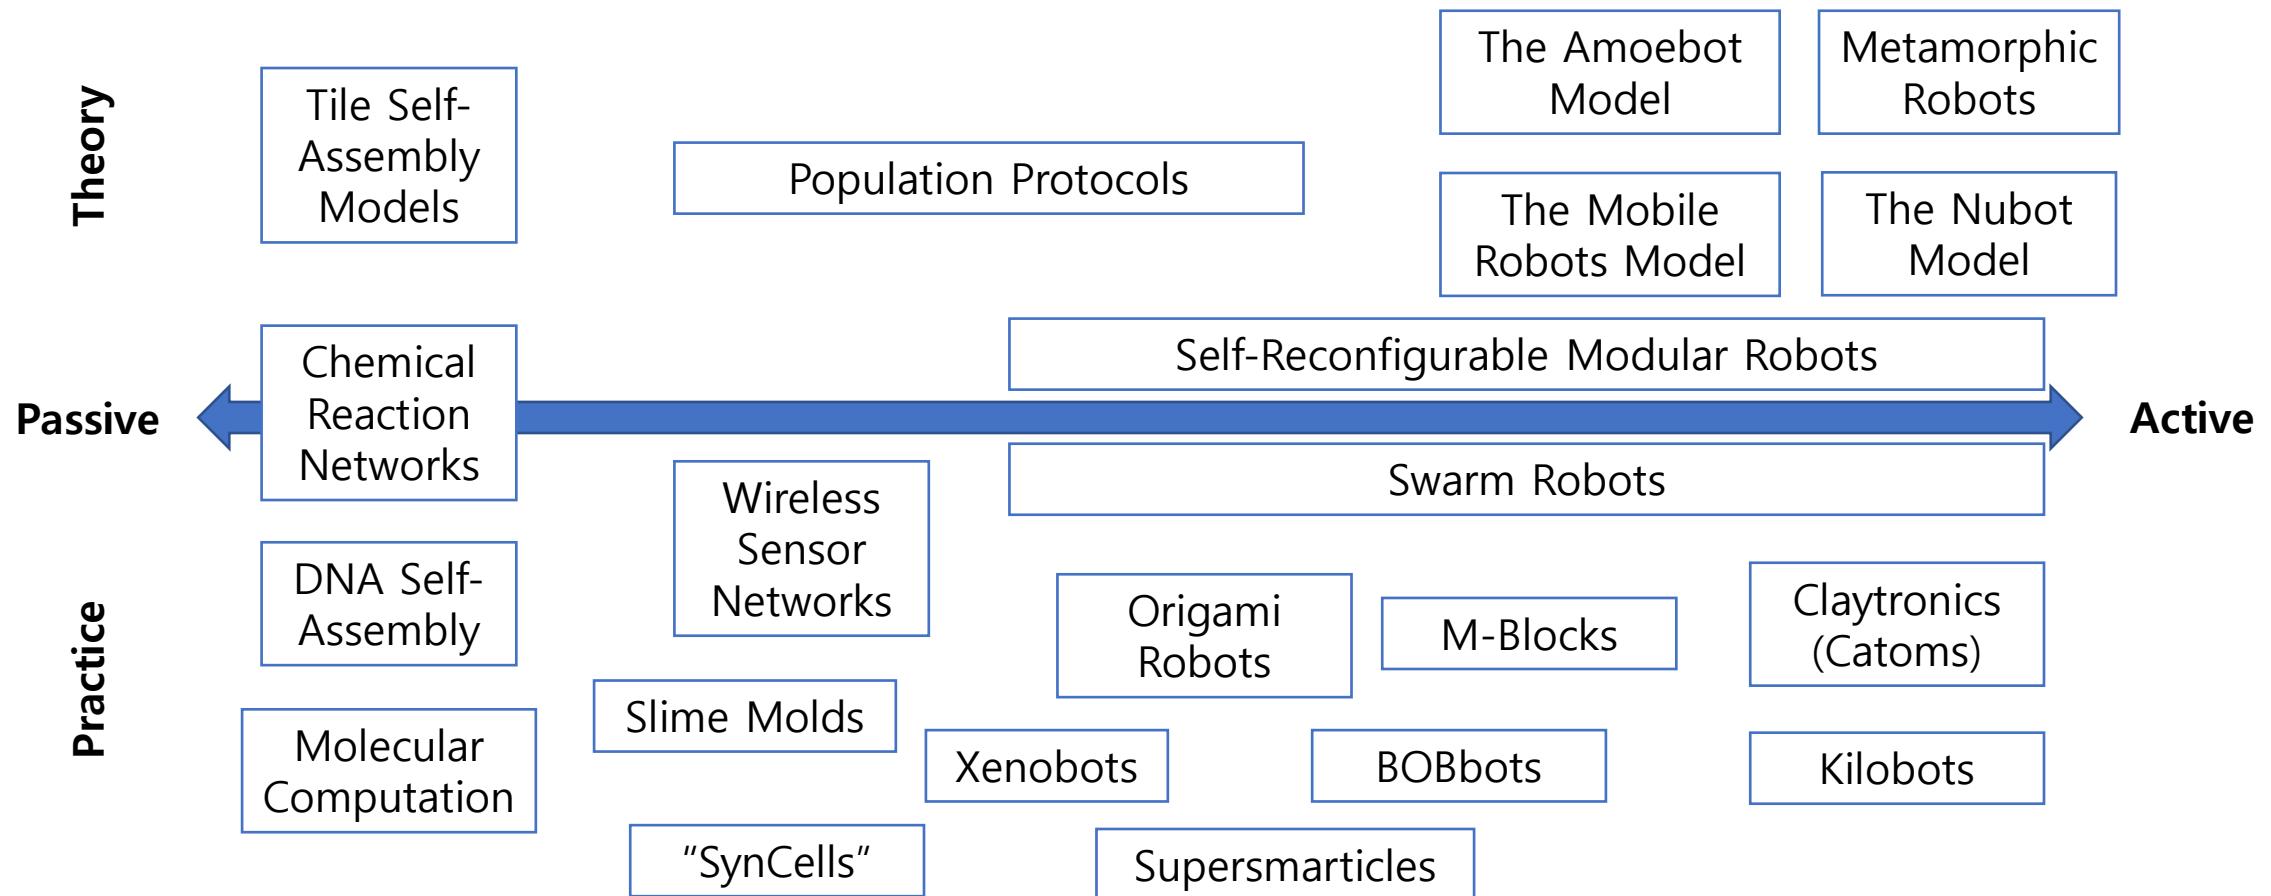


Most programmable matter and modular robotic systems assume:

- Modest [compute](#) resources.
- Strictly local [sensing](#) and [communication](#) (e.g., 1-neighborhood).
- Limited (e.g., constant-size) or no [persistent memory](#).
- Local, rudimentary [movement](#).

# Existing Research on Programmable Matter

Programmable matter systems can be organized by their **degree of self-determination** in deciding and enacting local behaviors.



# Dissertation: Outline

---

This dissertation focuses on the algorithmic foundations of active programmable matter.

Three main research questions:

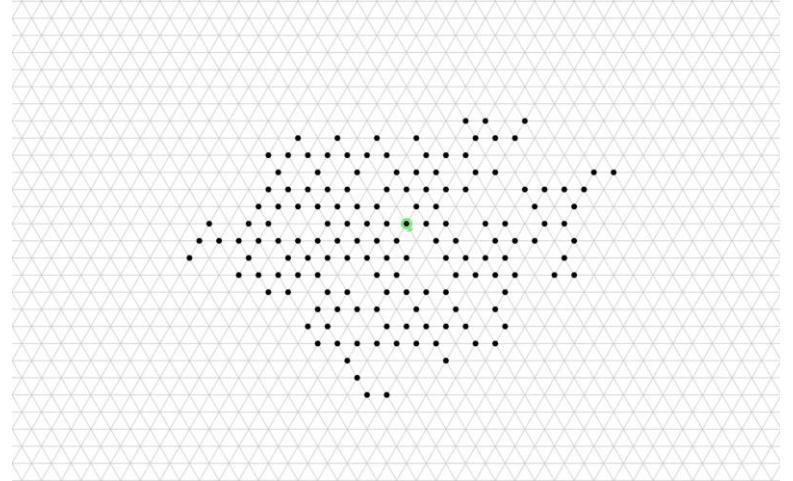
1. What are the minimum individual capabilities necessary to achieve system behavior  $X$ ?
2. How can existing algorithms be enhanced to capture more realistic assumptions?
3. How can digital algorithms be translated for simple, analog (passive) systems?

# Dissertation: The Big Picture

## 1. The Amoebot Model and its Enhancements



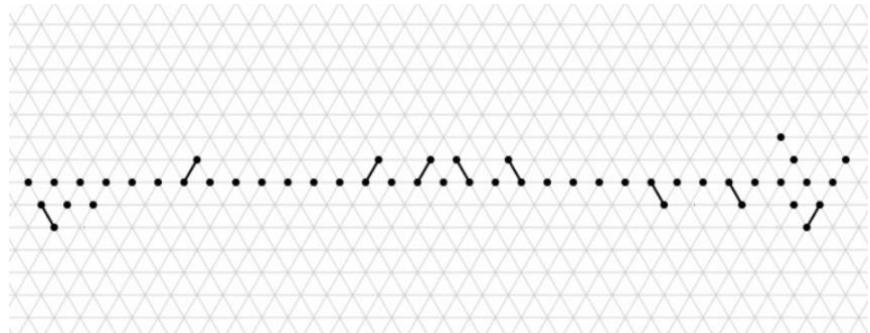
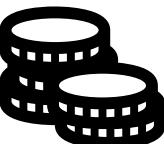
## 2. Stateful Distributed Algorithms



## 4. Swarm Robotics and Granular Active Matter



## 3. Stochastic Distributed Algorithms



# Part I

## Stateful Distributed Algorithms for Programmable Matter

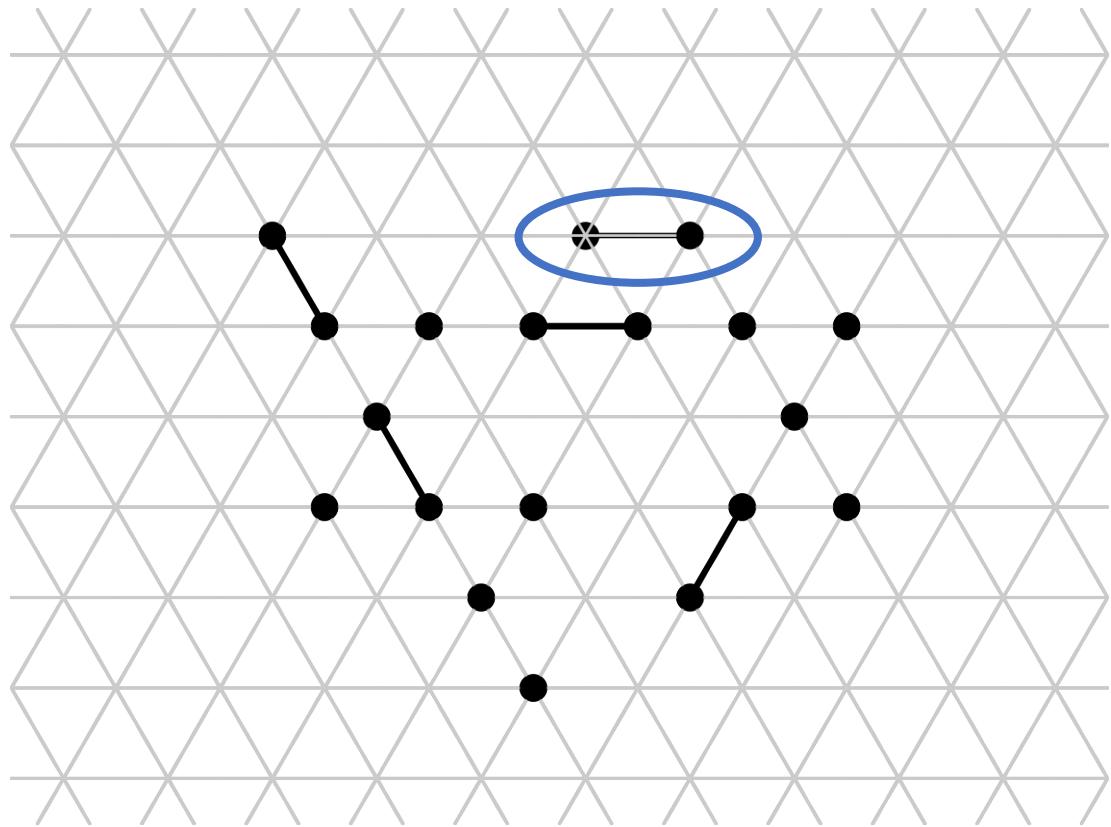
What are the minimum individual capabilities necessary to achieve system behavior X?

# The Amoebot Model (2014–2021)

---

The **amoebot model** is an abstraction of programmable matter.

- Space: triangular lattice  $G_\Delta$ .
- Amoebots can be **contracted** (one node) or **expanded** (two adjacent nodes).
- Amoebots are **anonymous**, have only **constant-size memories**, communicate with **immediate neighbors**, and have **no global compass**.
- Self-actuated movements via **expansions**, **contractions**, and **handovers**.
- **Sequential, weakly fair adversary**: one amoebot acts per time, every amoebot acts infinitely often.

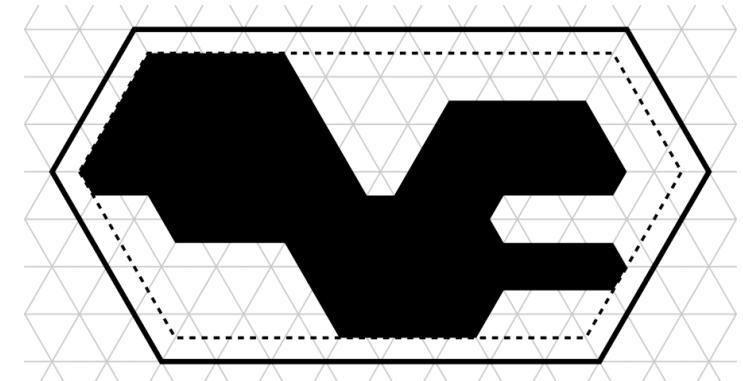
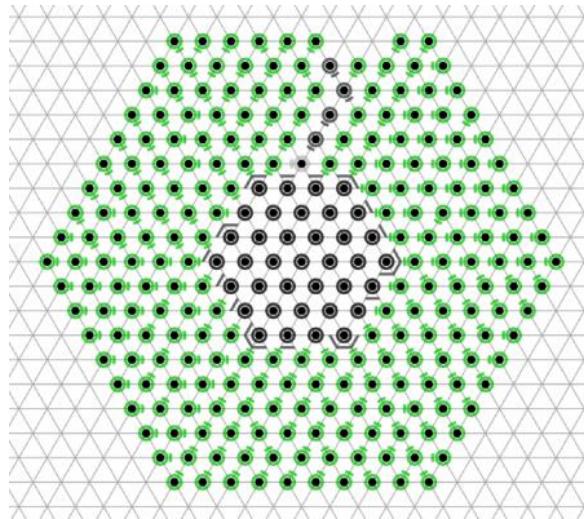
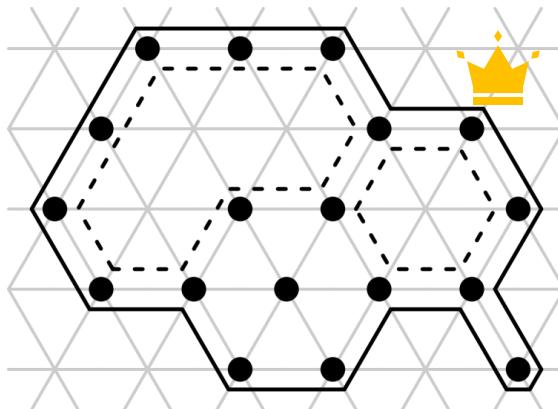


# Stateful Distributed Algorithms

---

With constant-size memory, communication between neighbors, and local movements, amoebot systems can solve:

1. **Leader Election.** A [unique](#) amoebot must [irreversibly](#) declare itself the system's [leader](#).
2. **Object Coating.** The system must reconfigure into even [layers](#) coating a given [object](#).
3. **Convex Hull Formation.** The system must reconfigure as the [convex hull](#) of a given [object](#), enclosing it with the [minimum number](#) of amoebots.

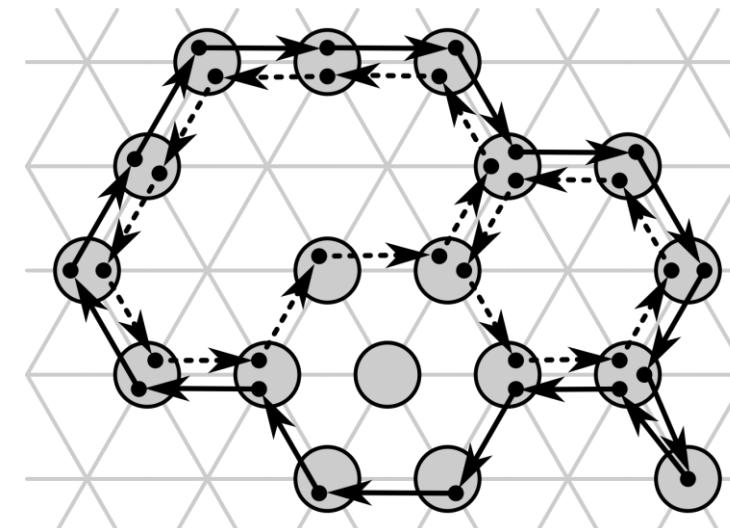
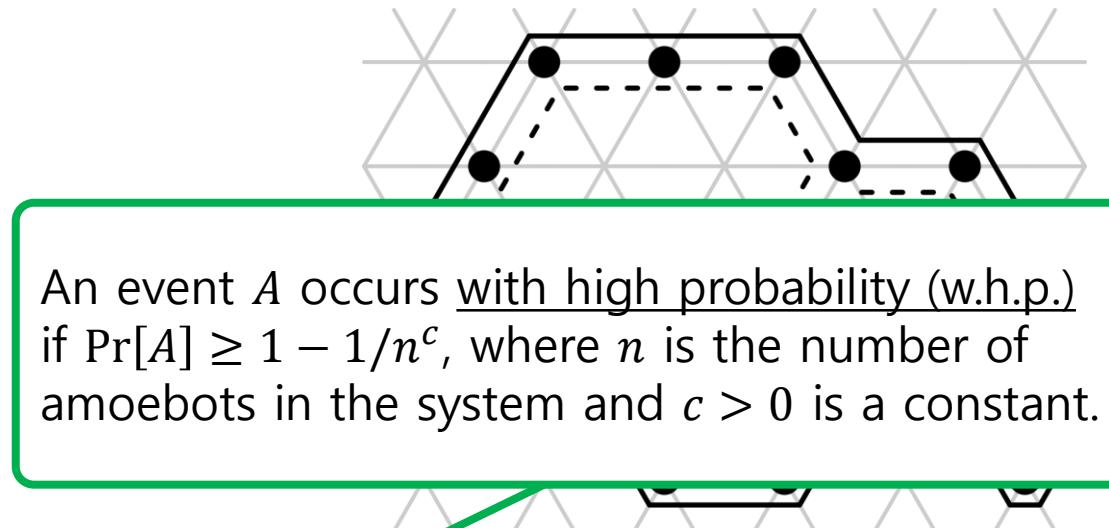


# Leader Election

Problem: A **unique** amoebot must **irreversibly** declare itself the system's **leader**.

Motivation: Leader election is well-studied in distributed computing. Can help coordinate the system for more complex behaviors (e.g., shape formation, object coating, etc.).

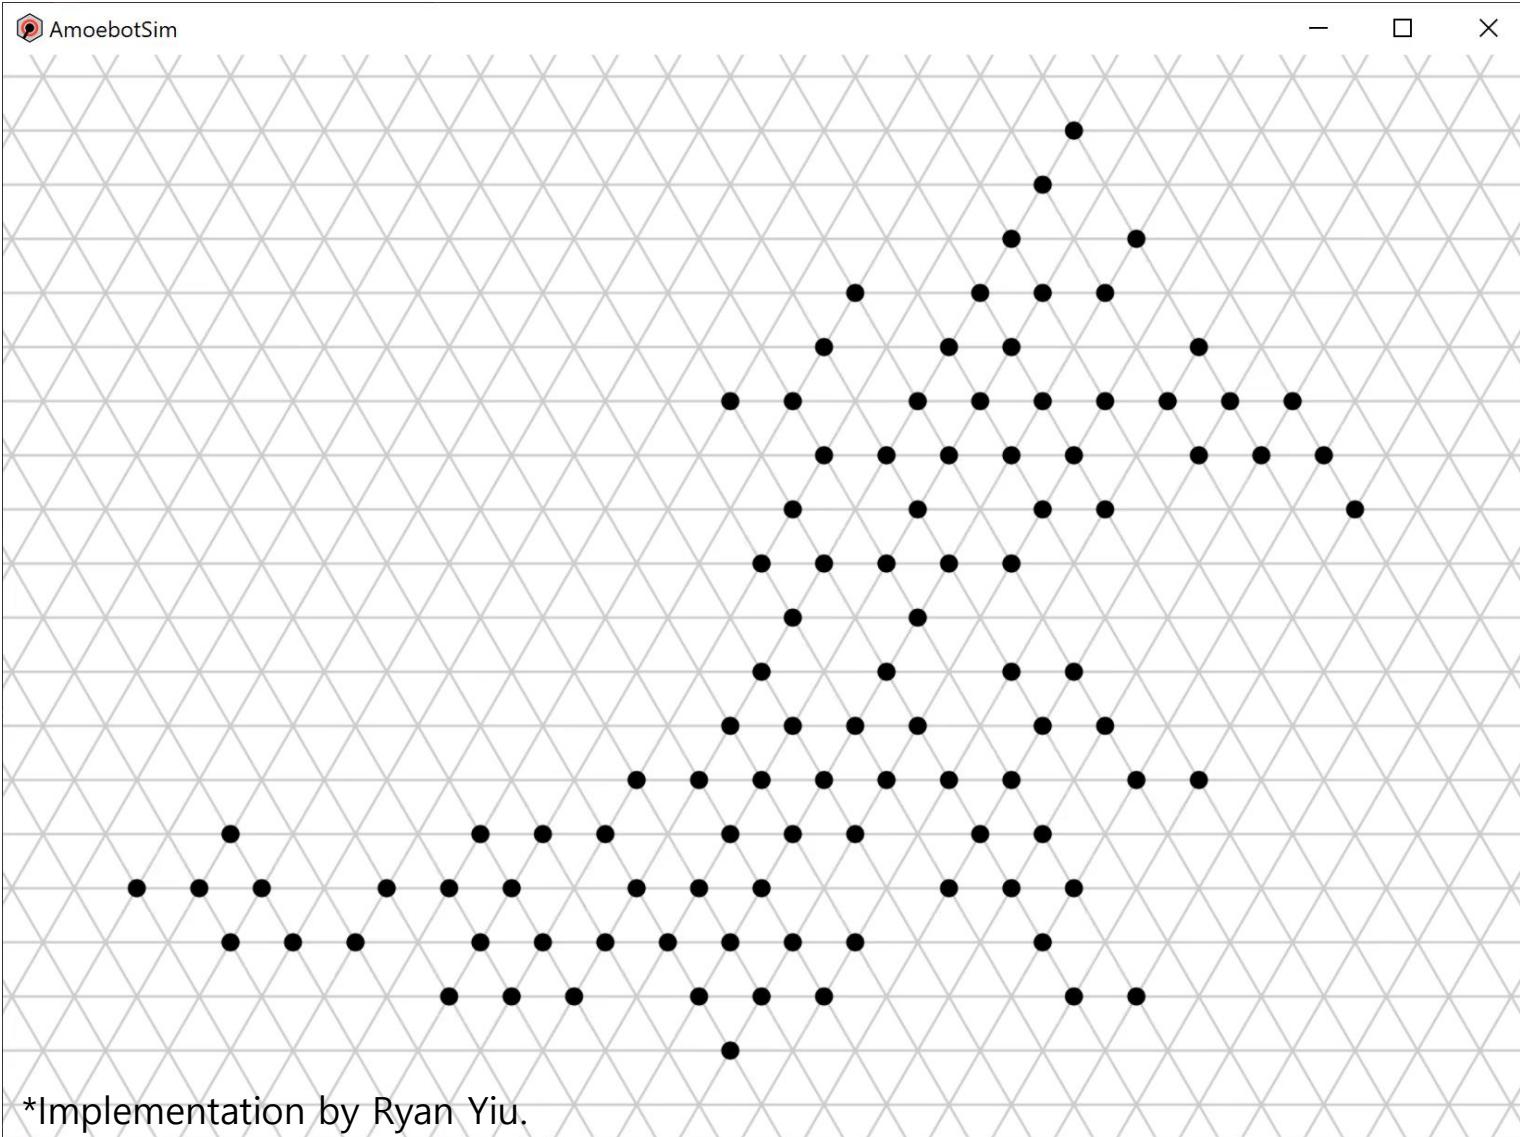
Key Idea: Each amoebot gets a **random value**. The **largest value** on the **outer boundary** wins.



**Theorem**. The Improved-Leader-Election algorithm solves the leader election problem in  $\mathcal{O}(L)$  rounds w.h.p., where  $L$  is the length of the outer boundary.

# Leader Election

---



# Leader Election

---

Problem: A unique amoebot must irreversibly declare itself the system's leader.

**Theorem.** The Improved-Leader-Election algorithm solves the leader election problem in  $\mathcal{O}(L)$  rounds w.h.p., where  $L$  is the length of the outer boundary.

Our algorithm inspired significant follow-up work:

Algorithm	Det.	Weak Sched.	Allows Holes	Removes Chirality	Static	Leaders Elected	Runtime
Leader-Election [59]	No	No	Yes	No	Yes	1	$\mathcal{O}(L^*)$ exp.
Improved-Leader-Election	No	No	Yes	No	Yes	1, whp.	$\mathcal{O}(L)$ whp.
Di Luna et al. [64, 65]	Yes	Yes	No	Yes	Yes	$k \leq 3$	$\mathcal{O}(n^2)$
Gastineau et al. [91]	Yes	No	No	No	Yes	1	$\mathcal{O}(n)$
Bazzi and Briones [21]	Yes	Yes	Yes	No	Yes	$k \leq 6$	$\mathcal{O}(n^2)$
Emek et al. [77]	Yes	No	Yes	Yes	No	1	$\mathcal{O}(Ln^2)$

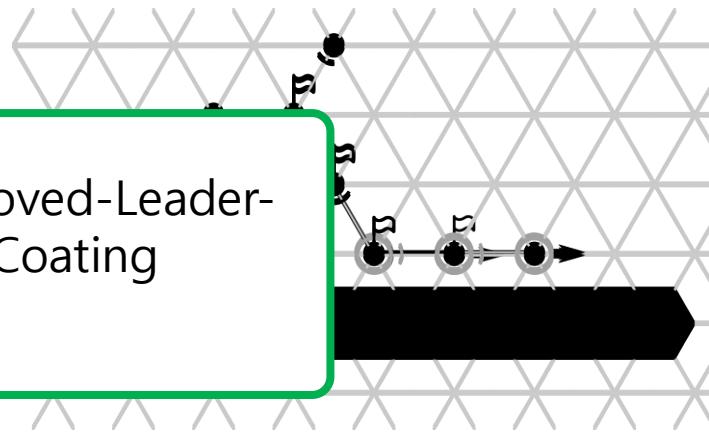
Interestingly, our Improved-Leader-Election algorithm remains state-of-the-art for settings with holes where amoebots cannot move and exactly one leader should be elected.

# Object Coating

Problem: The system must reconfigure into even layers coating a given object.

Motivation: Smart paint, distributed sensor networks, and shape formation via reverse-molds.

Key Idea: Coat the first layer by following the object's surface. Elect a leader to mark the start and end of higher layers. [DGRSS 2017].

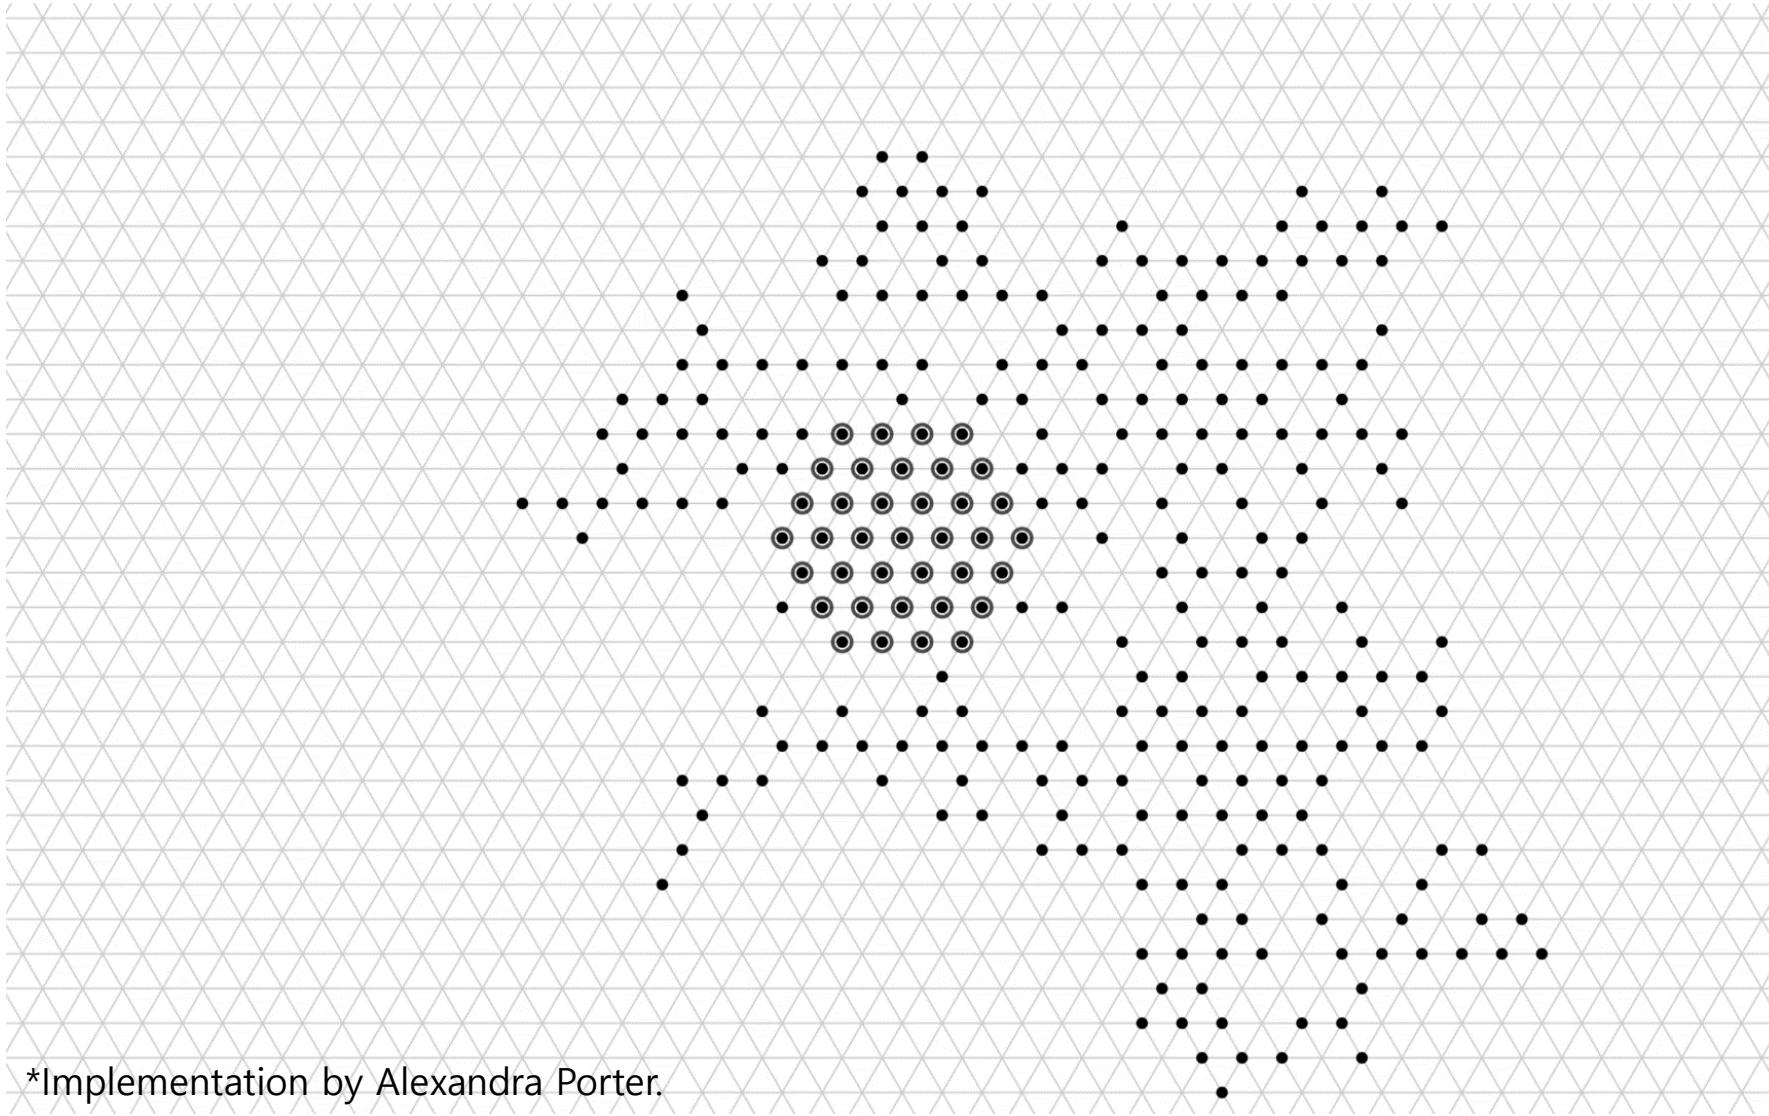


This "w.h.p." is inherited from Improved-Leader-Election. The rest of the Universal-Coating algorithm is deterministic.

**Theorem**. The Universal-Coating algorithm solves the object coating problem in  $\mathcal{O}(n)$  rounds w.h.p., where  $n$  is the number of amoebots in the system. This runtime is worst-case asymptotically optimal — no local-control algorithm can do any better, in the worst case.

# Object Coating

---



\*Implementation by Alexandra Porter.

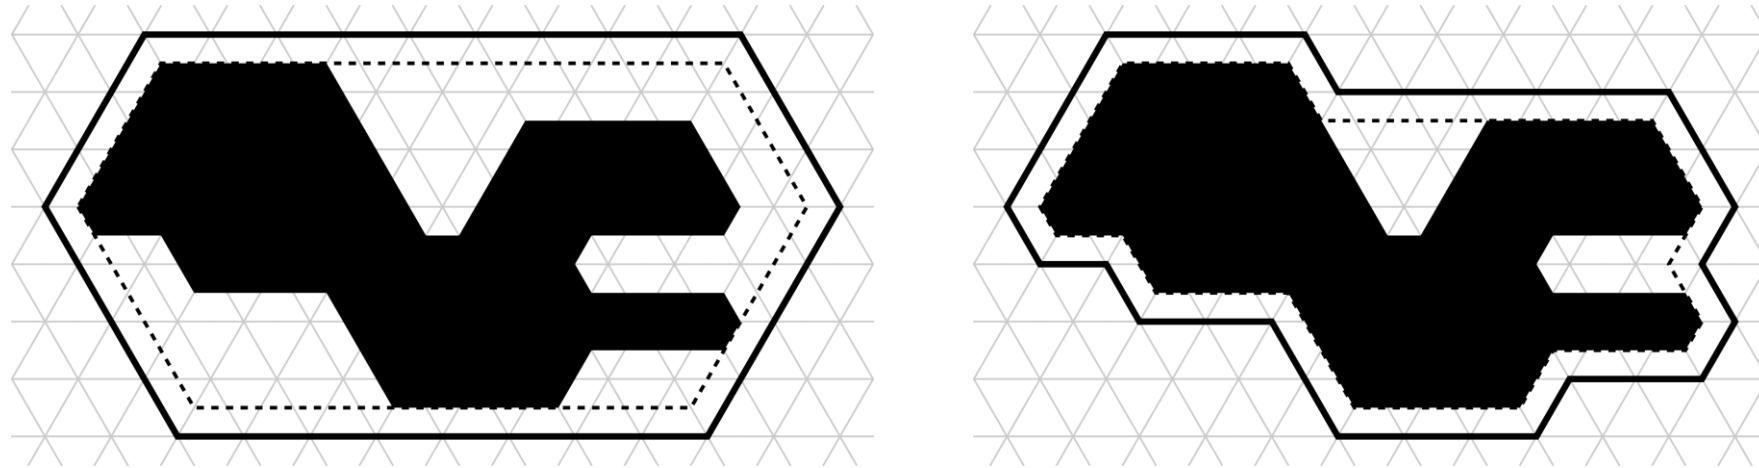
# Convex Hull Formation

---

Problem: The system must reconfigure as the [convex hull](#) of a given [object](#), enclosing it with the [minimum number](#) of amoebots.

Motivation: Collective transport, isolating hazardous materials, macrophage-like engulfing.

In our discrete setting of the triangular lattice, we consider [restricted-orientation convex hulls](#).

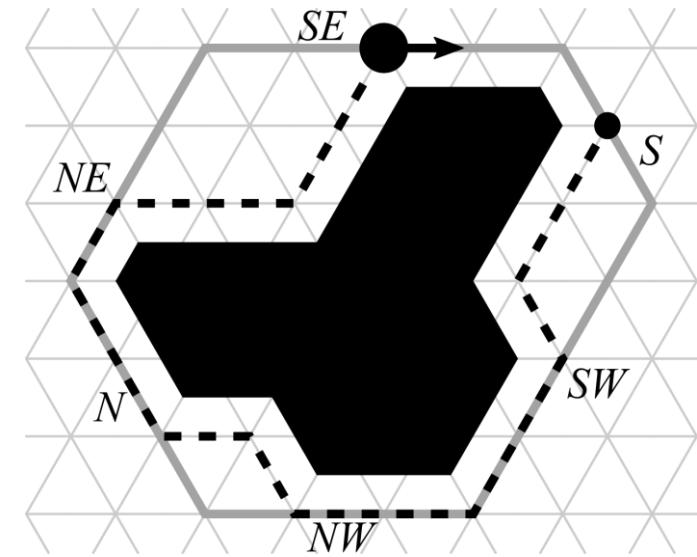
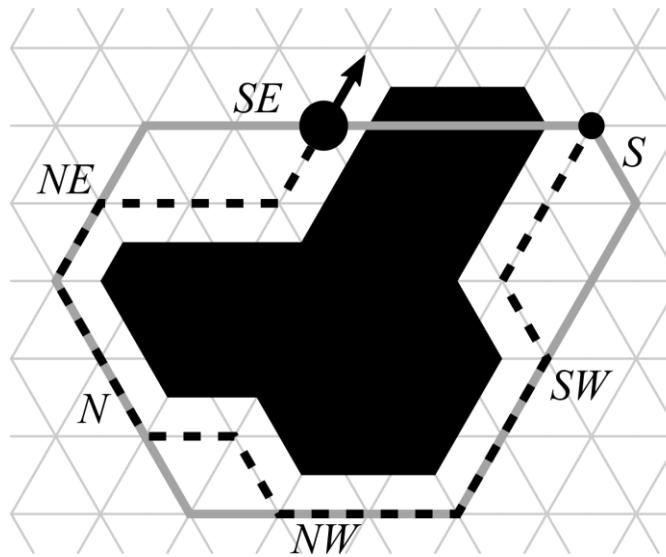
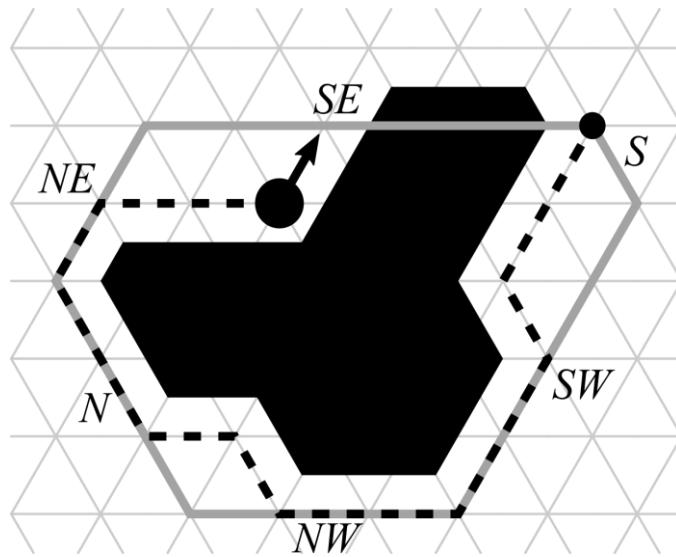


Ours is the [first distributed algorithm](#) to compute restricted-orientation convex hulls [without global orientation or coordinates](#) and when limited to [constant-size memory](#).

# Convex Hull Formation

Problem: The system must reconfigure as the **convex hull** of a given **object**, enclosing it with the **minimum number** of amoebots.

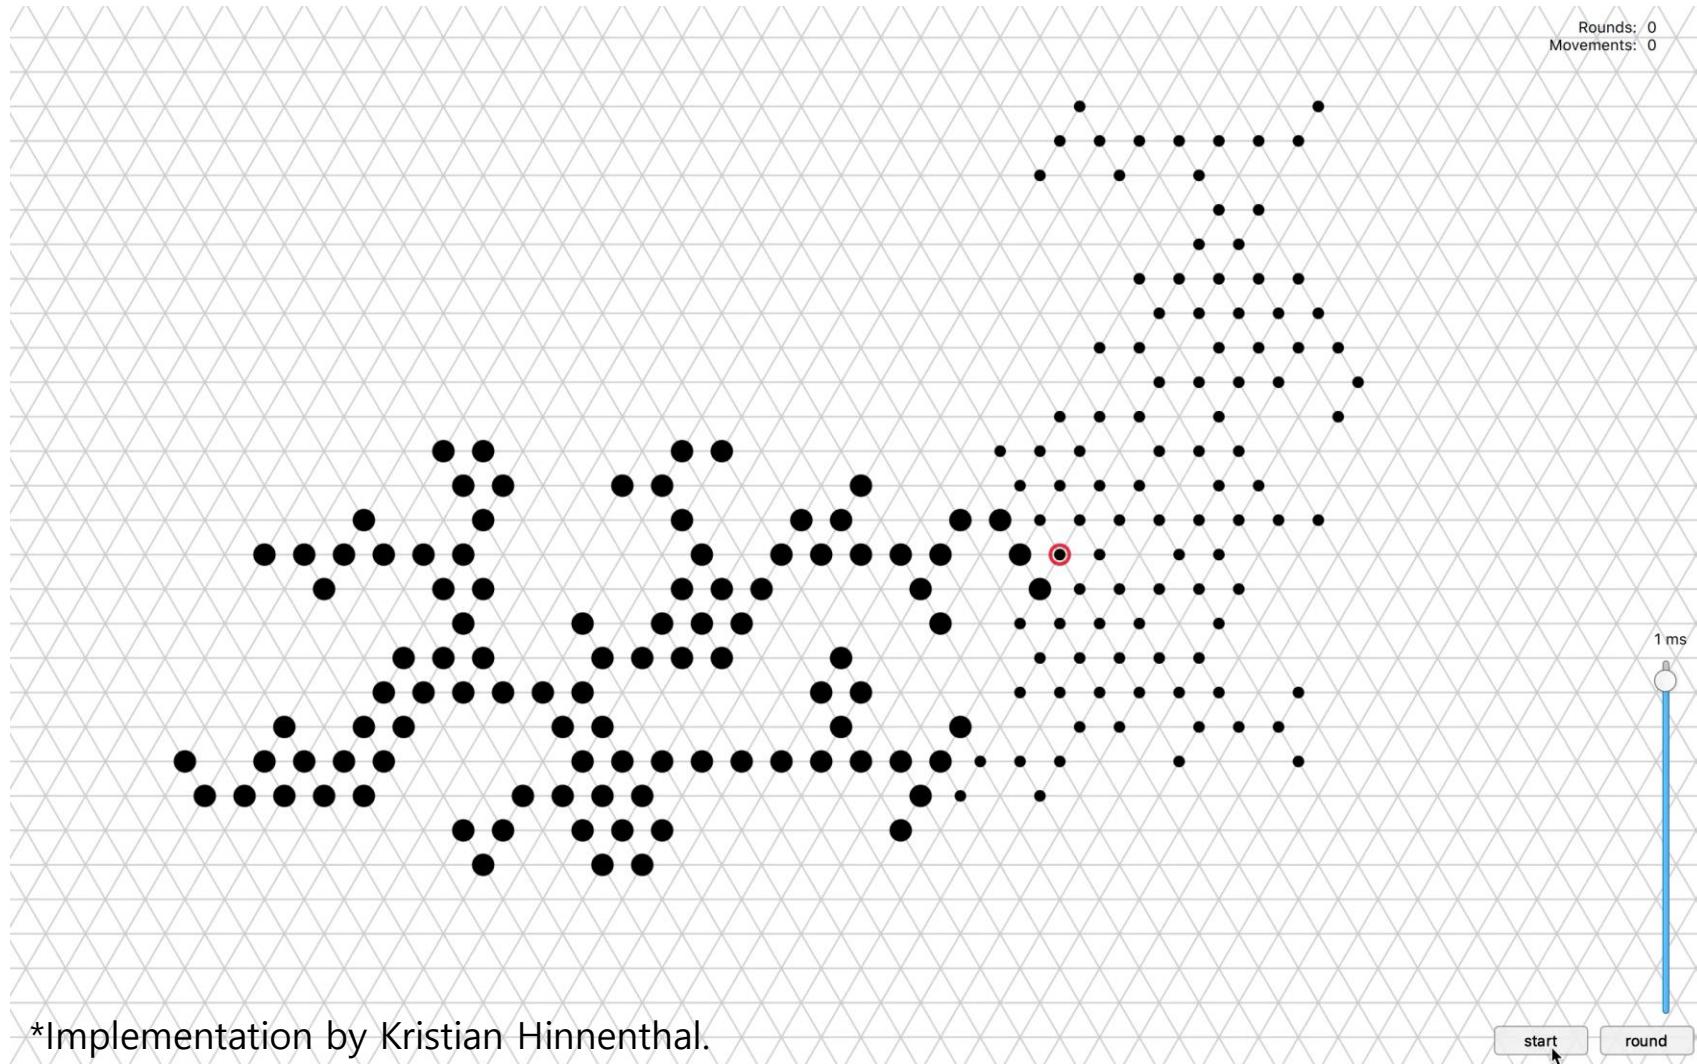
Key Idea: Use a leader to explore the object, keeping track of its **distances** to each of the six **half-planes** forming the convex hull. Once determined, simply follow the convex hull.



But distances are **too big for constant-size memory**! So we use the rest of the amoebot system as the leader's **distributed memory**.

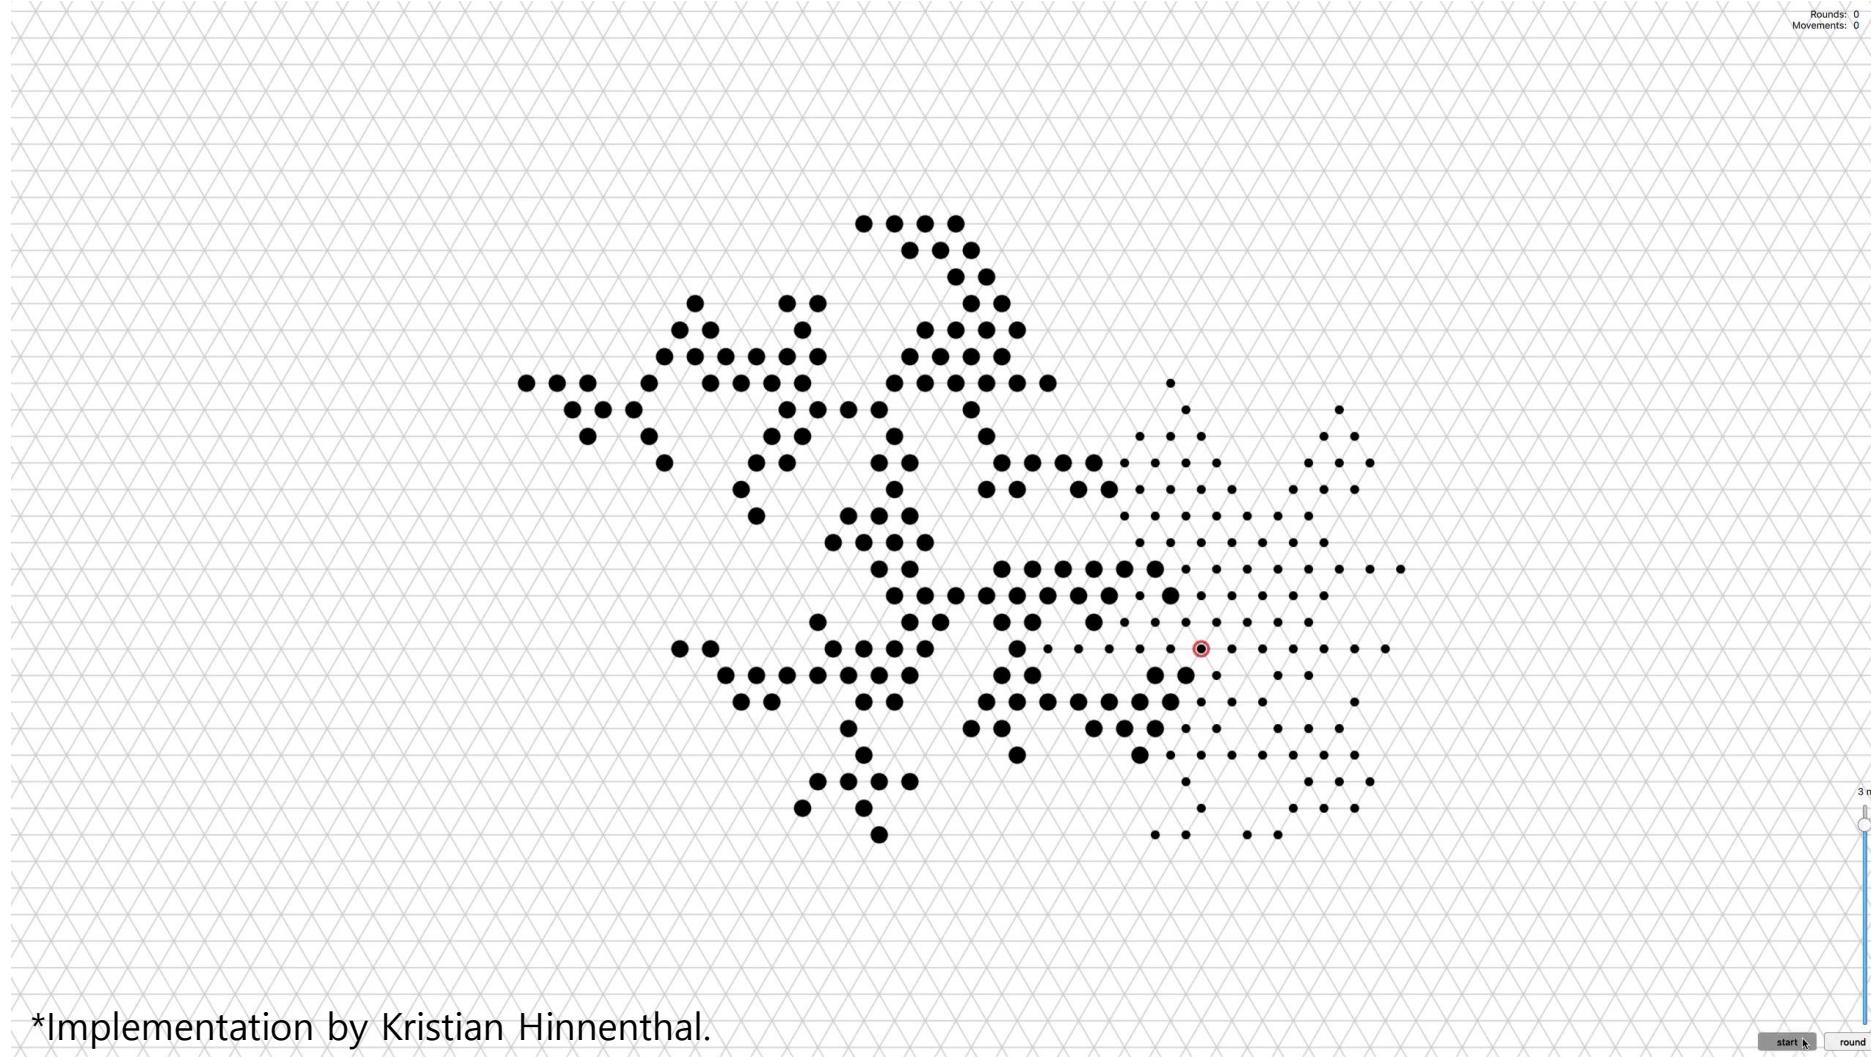
# Convex Hull Formation

---



# Convex Hull Formation

---



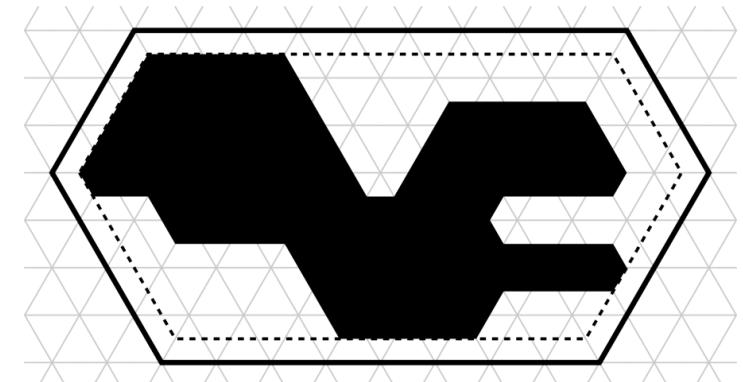
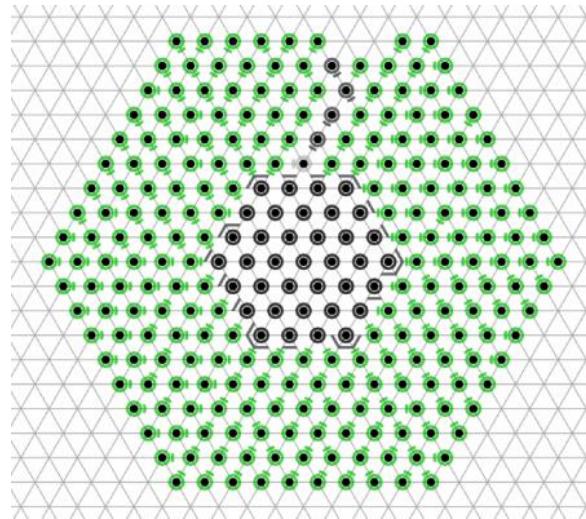
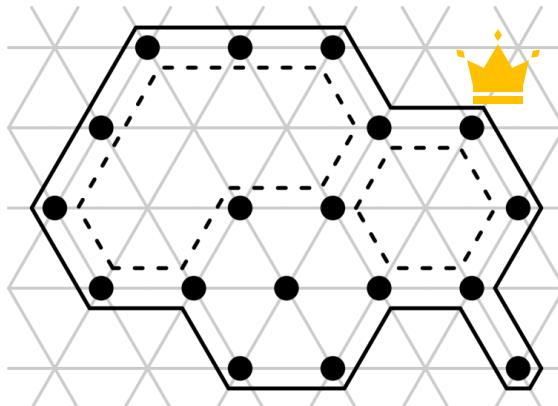
\*Implementation by Kristian Hinnenthal.

# Part I

## Stateful Distributed Algorithms for Programmable Matter

What are the **minimum individual capabilities** necessary to achieve **system behavior X**?

With **constant-size memory**, **neighbor-to-neighbor communication**, and **local movements**, a system can collectively achieve leader election, object coating, and convex hull formation.



## Part II

### The Amoebot Model and its Enhancements

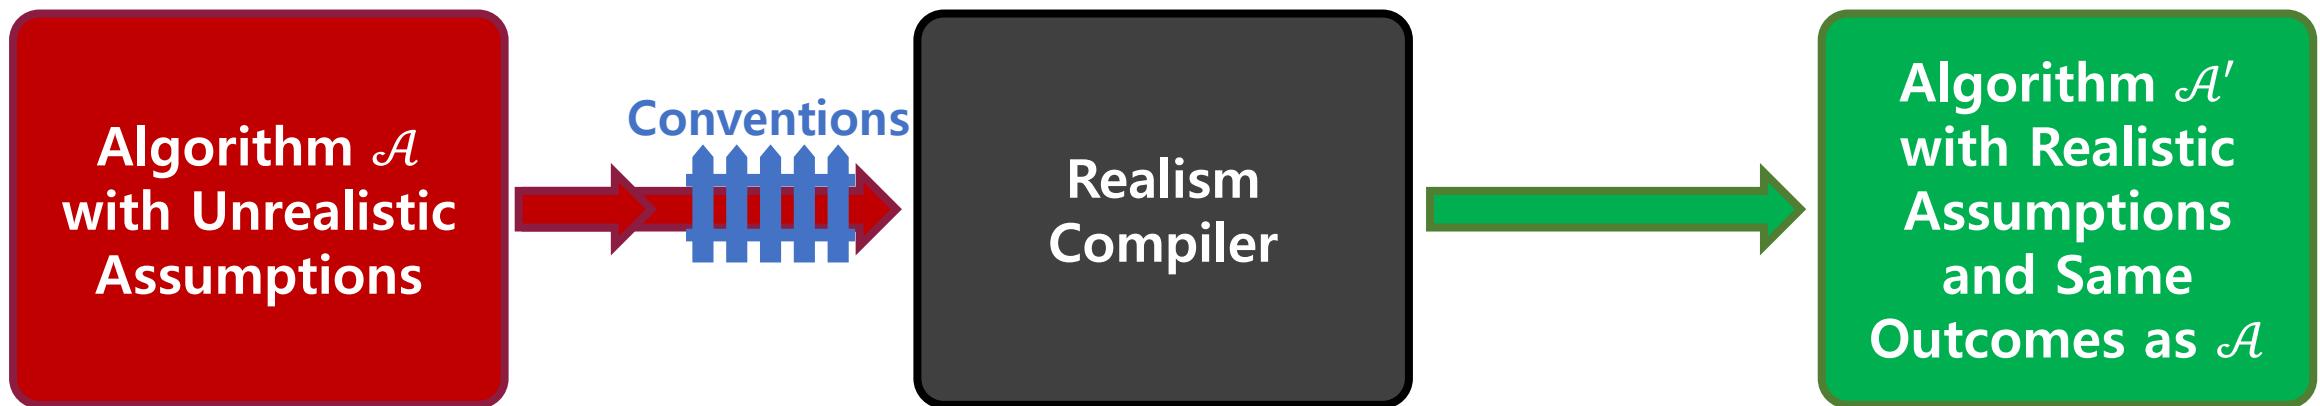
How can **existing algorithms** be enhanced to capture more **realistic assumptions**?

# Enhancing the Amoebot Model

The amoebot model and its algorithms do not account for energy costs of the amoebots' actions (**energy-agnostic**) and assume only one amoebot is active at a time (**sequential**).

Real programmable matter systems are **energy-constrained** and **concurrent**.

At a high level, what we'd like is the following:



This is **too optimistic** and may be **impossible** to guarantee in general, so instead we only consider algorithms  $\mathcal{A}$  that obey certain **conventions**.

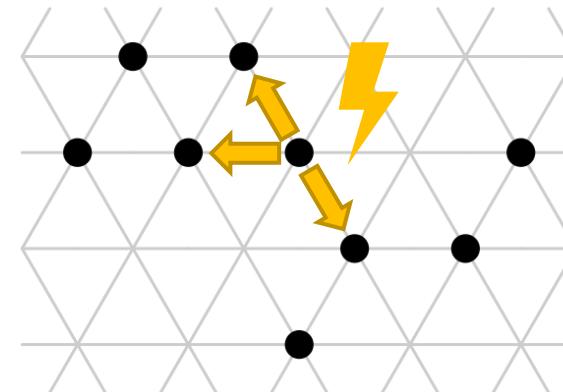
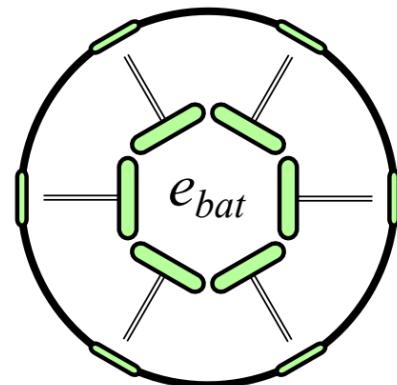
# Energy Distribution

---

Goal: Model energy harvesting, distribution, and usage. Ensure all amoebots eventually get the energy they need to run some algorithm  $\mathcal{A}$ .

## Model Extensions

- Each amoebot  $A$  has a constant-size **battery**  $A.e_{bat}$ .
- Amoebots with access to an **external energy source** can directly harvest energy.
- Amoebots can **transfer** a fixed amount of energy per time to their **neighbors** without loss.



# Energy Distribution

Goal: Model energy harvesting, distribution, and usage. Ensure all amoebots eventually get the energy they need to run some algorithm  $\mathcal{A}$ .

We developed the **Energy-Sharing** algorithm as an asymptotically optimal mechanism for **distributing energy** to all amoebots in a system.

We then developed the **Forest-Prune-Repair** algorithm as a mechanism for maintaining an underlying **spanning forest** structure as amoebots **move**.

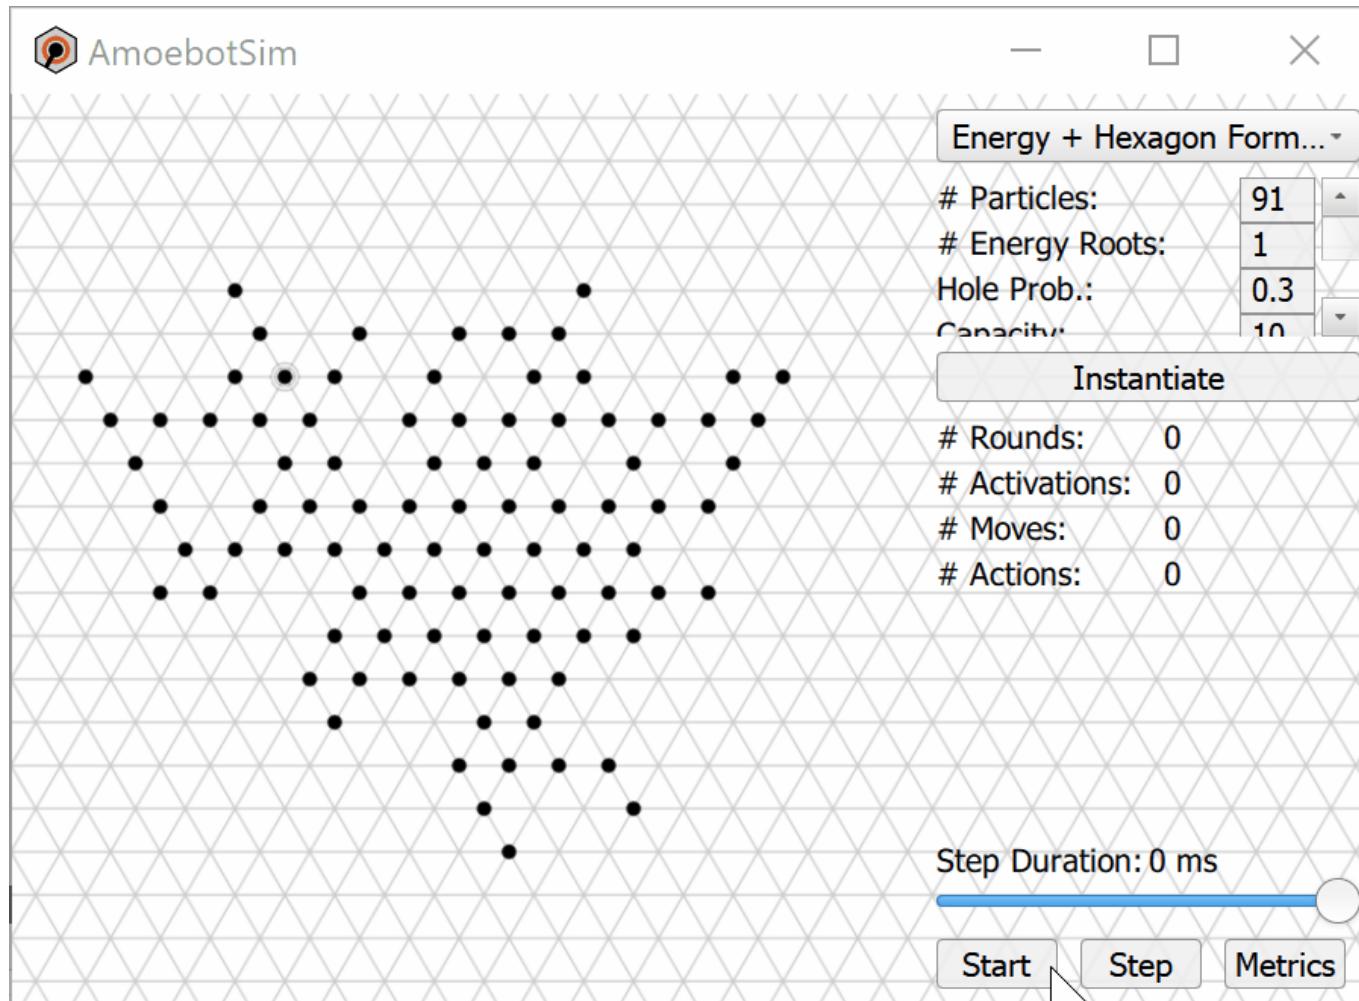


\*Joint work with Jamison Weber.



# Energy-Constrained Shape Formation

Energy-Sharing + Forest-Prune-Repair composed with **Hexagon-Formation**:



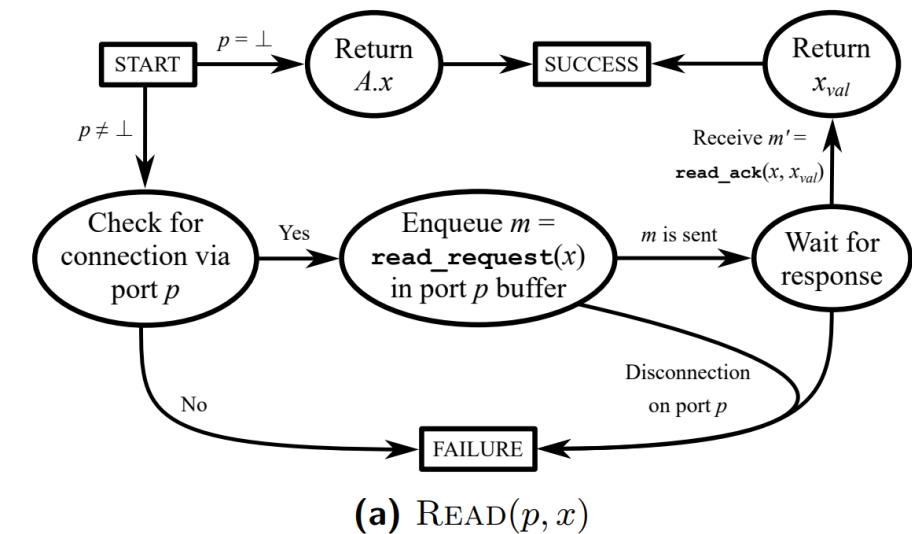
# The Canonical Amoebot Model

Goal: Study amoebot algorithms where many amoebots are simultaneously active.

Generalizes the amoebot model by partitioning amoebot functionality into:

- A higher-level **application layer** where algorithms are defined in terms of [operations](#).
- A lower-level **system layer** that executes an amoebot's operations via [message passing](#).

Operation	Return Value on Success
$\text{CONNECTED}(p)$	TRUE iff a neighboring amoebot is connected via port $p$
$\text{READ}(p, x)$	The value of $x$ in the public memory of this amoebot if $p = \perp$ or of the neighbor incident to port $p$ otherwise
$\text{WRITE}(p, x, x_{val})$	Confirmation that the value of $x$ was updated to $x_{val}$ in the public memory of this amoebot if $p = \perp$ or of the neighbor incident to port $p$ otherwise
$\text{CONTRACT}(v)$	Confirmation of the contraction out of node $v \in \{\text{HEAD}, \text{TAIL}\}$
$\text{EXPAND}(p)$	Confirmation of the expansion into the node incident to port $p$
$\text{PULL}(p)$	Confirmation of the pull handover with the neighbor incident to port $p$
$\text{PUSH}(p)$	Confirmation of the push handover with the neighbor incident to port $p$
$\text{LOCK}()$	Local identifiers of the amoebots that were successfully locked
$\text{UNLOCK}(\mathcal{L})$	Confirmation that the amoebots of $\mathcal{L}$ were unlocked



# The Canonical Amoebot Model

---

Algorithms in the canonical amoebot model are specified in terms of **actions**:

$$\langle \text{label} \rangle : \langle \text{guard} \rangle \rightarrow \langle \text{operations} \rangle$$

- *label* specifies the action's name.
- *guard* is a Boolean predicate determining whether this action is currently **enabled**.
- *operations* specifies the computation and sequence of operations to perform if enacted.

Example from Hexagon-Formation:

$$\underline{\alpha_2 : (A.\text{state} = \text{IDLE}) \wedge (\exists B \in N(A) : B.\text{state} \in \{\text{FOLLOWER}, \text{ROOT}\})} \rightarrow$$

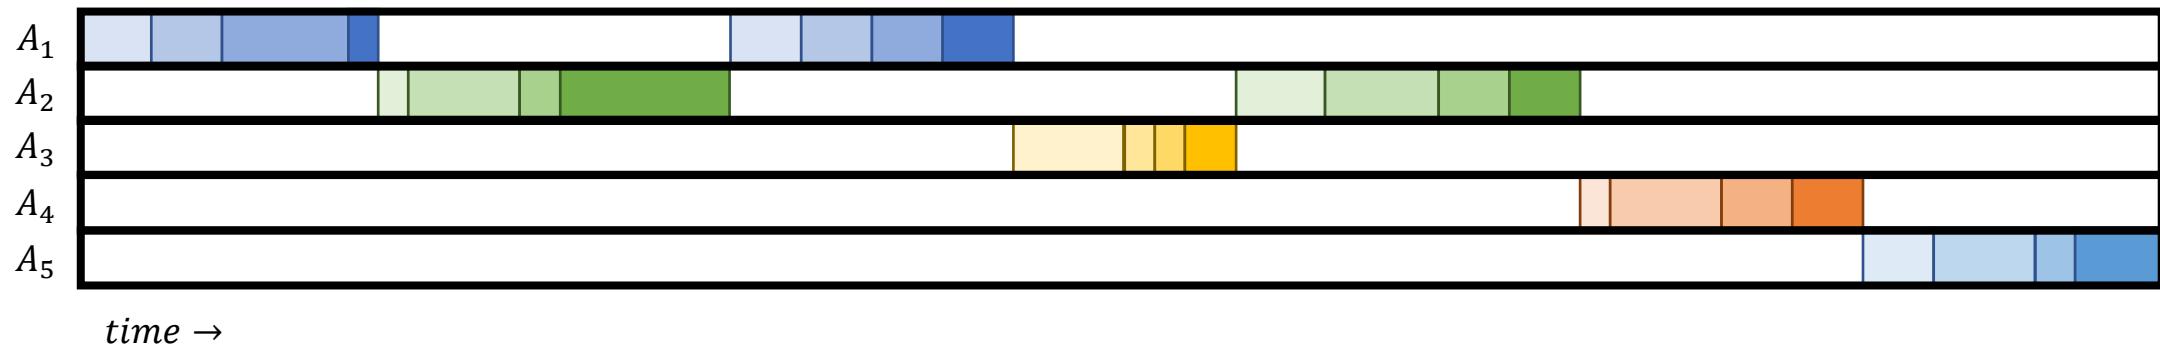
Find a port  $p$  for which  $\text{CONNECTED}(p) = \text{TRUE}$  and  $\text{READ}(p, \text{state}) \in \{\text{FOLLOWER}, \text{ROOT}\}$ .  
 $\text{WRITE}(\perp, \text{parent}, p)$ .  
 $\text{WRITE}(\perp, \text{state}, \text{FOLLOWER})$ .

# The Canonical Amoebot Model

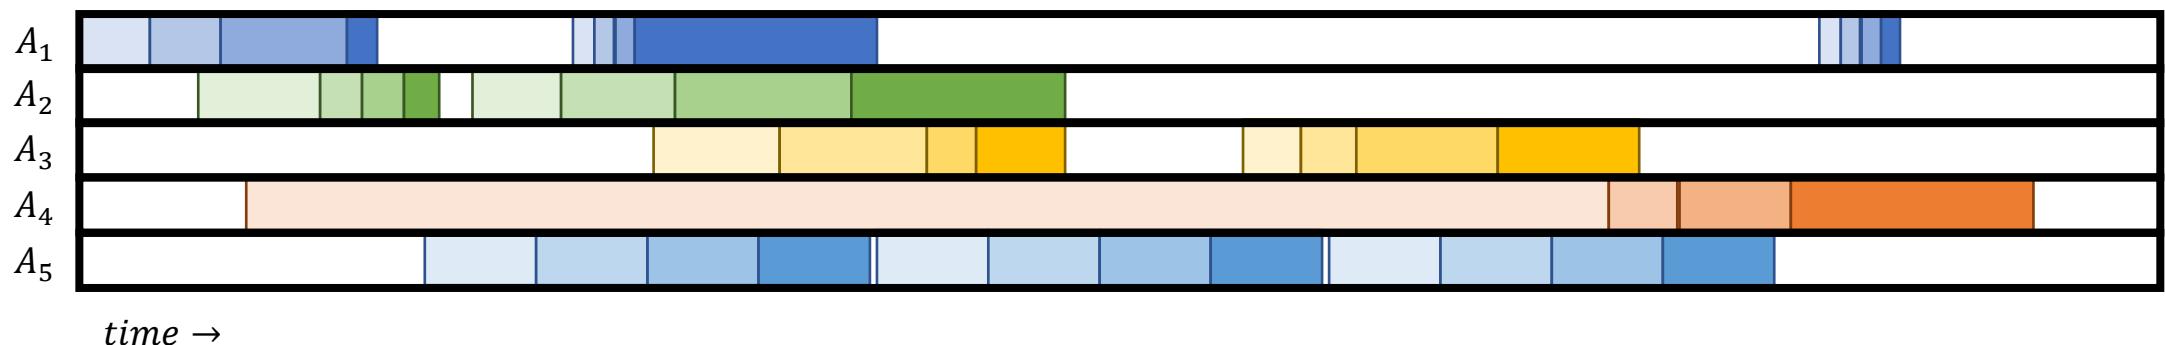
---

We use an **adversary** to model timing and progress. Two primary levels of **concurrency**:

**Sequential.** At most one active amoebot per time.



**Asynchronous.** Arbitrary sets of amoebots can be simultaneously active.



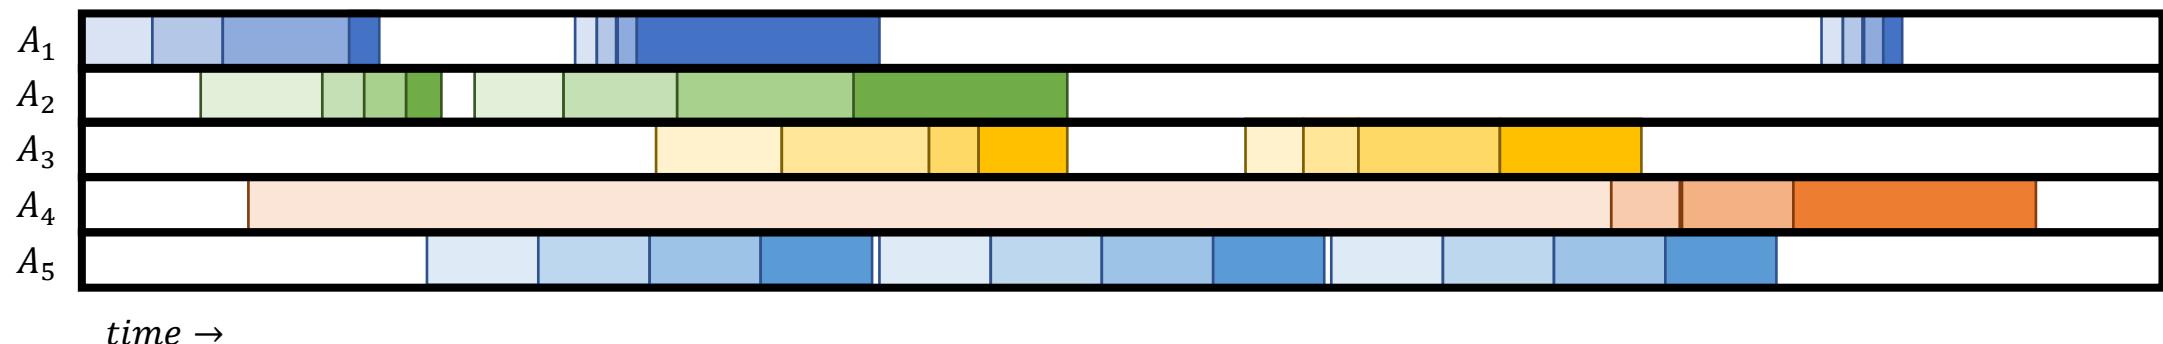
# The Canonical Amoebot Model

---

An **unfair** adversary can activate any amoebot with an enabled action.

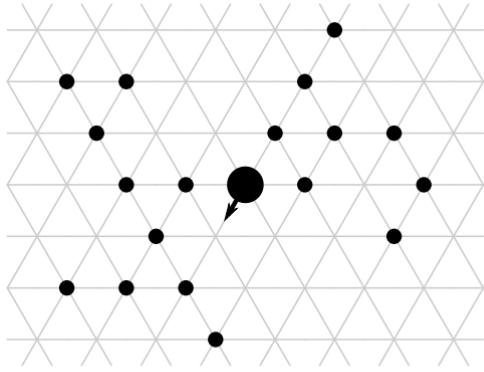
The rest of this talk will primarily focus on [unfair asynchronous adversaries](#), the most general of all adversarial activation models.

Informally: the adversary can activate any amoebot with something to do whenever it wants to.

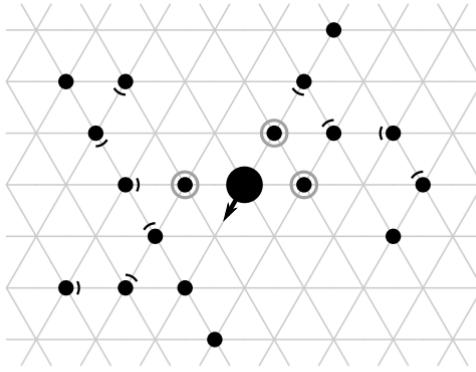


# Asynchronous Hexagon Formation

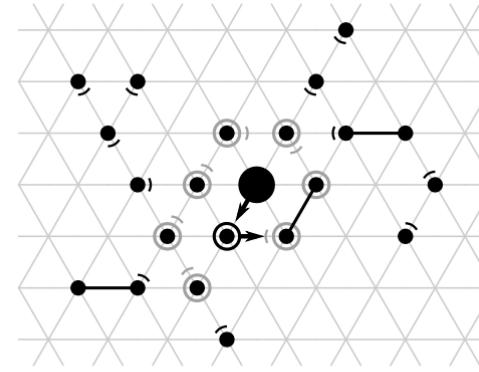
Problem: Reconfigure any connected amoebot system as a regular hexagon, assuming there is a unique seed amoebot initially in the system.



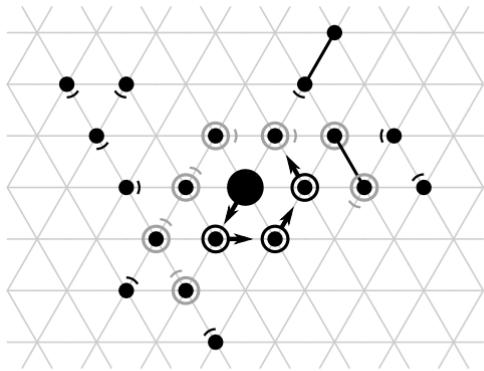
(a)



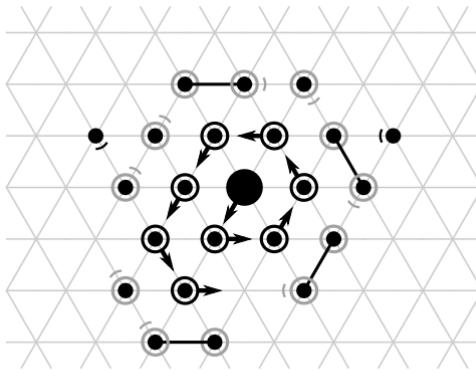
(b)



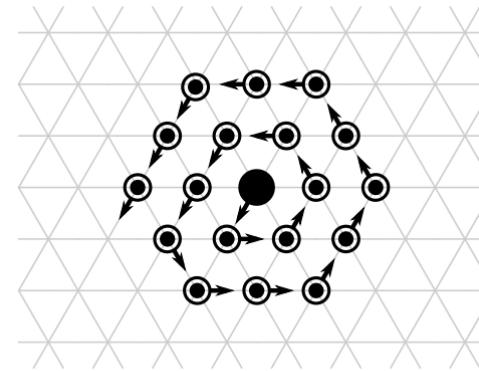
(c)



(d)



(e)



(f)

# Asynchronous Hexagon Formation

---

We formulate a Hexagon-Formation algorithm in terms of actions based on [DGRSS 2015].

## Algorithm 1 Hexagon-Formation for Amoobot $A$

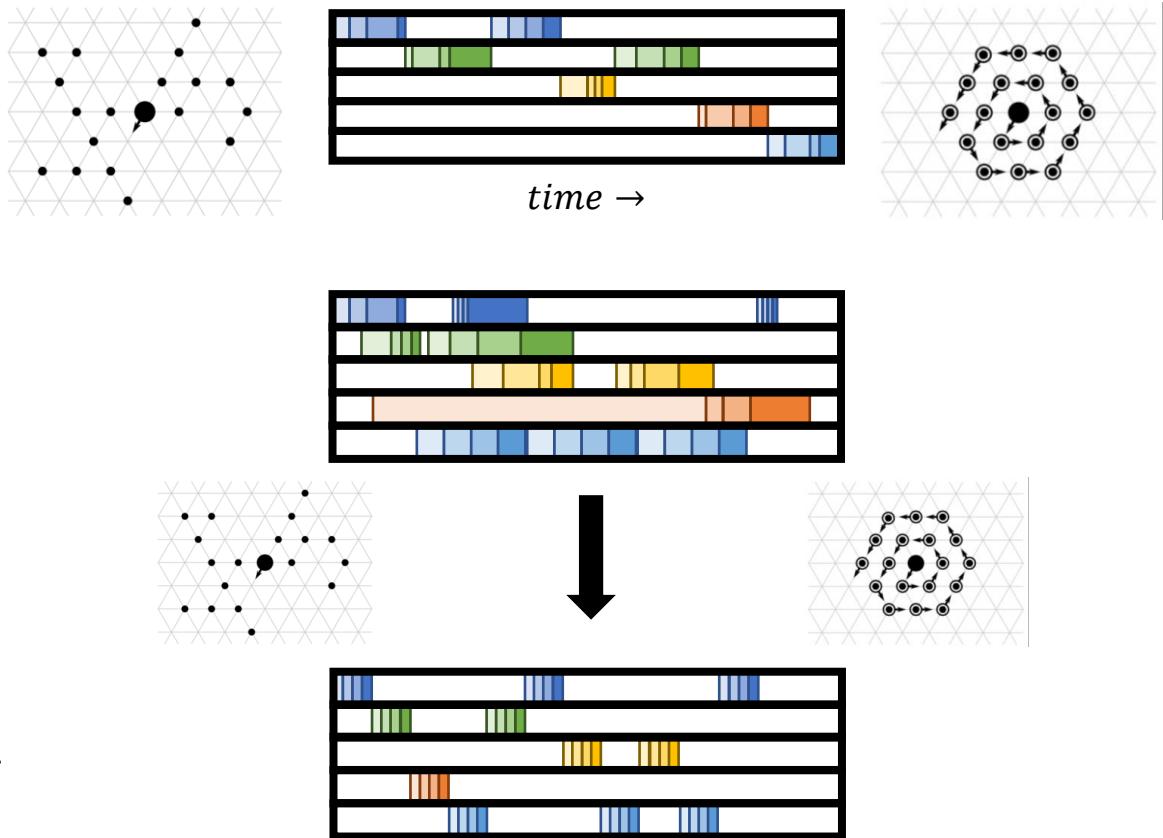
- 
- 1:  $\alpha_1 : (A.\text{state} \in \{\text{IDLE}, \text{FOLLOWER}\}) \wedge (\exists B \in N(A) : B.\text{state} \in \{\text{SEED}, \text{RETIRED}\}) \rightarrow$
  - 2:    $\text{WRITE}(\perp, \text{parent}, \text{NULL}).$
  - 3:    $\text{WRITE}(\perp, \text{state}, \text{ROOT}).$
  - 4:    $\text{WRITE}(\perp, \text{dir}, \text{GETNEXTDIR(counter-clockwise)}).$  ▷ See Algorithm 2.
  - 5:  $\alpha_2 : (A.\text{state} = \text{IDLE}) \wedge (\exists B \in N(A) : B.\text{state} \in \{\text{FOLLOWER}, \text{ROOT}\}) \rightarrow$
  - 6:   Find a port  $p$  for which  $\text{CONNECTED}(p) = \text{TRUE}$  and  $\text{READ}(p, \text{state}) \in \{\text{FOLLOWER}, \text{ROOT}\}.$
  - 7:    $\text{WRITE}(\perp, \text{parent}, p).$
  - 8:    $\text{WRITE}(\perp, \text{state}, \text{FOLLOWER}).$
  - 9:  $\alpha_3 : (A.\text{shape} = \text{CONTRACTED}) \wedge (A.\text{state} = \text{ROOT}) \wedge (\forall B \in N(A) : B.\text{state} \neq \text{IDLE})$
  - 10:    $\wedge (\exists B \in N(A) : (B.\text{state} \in \{\text{SEED}, \text{RETIRED}\}) \wedge (B.\text{dir is connected to } A)) \rightarrow$
  - 11:    $\text{WRITE}(\perp, \text{dir}, \text{GETNEXTDIR(clockwise)}).$
  - 12:    $\text{WRITE}(\perp, \text{state}, \text{RETIRED}).$
  - 13:  $\alpha_4 : (A.\text{shape} = \text{CONTRACTED}) \wedge (A.\text{state} = \text{ROOT}) \wedge (\text{the node adjacent to } A.\text{dir is empty}) \rightarrow$
  - 14:    $\text{EXPAND}(A.\text{dir}).$
  - 15:  $\alpha_5 : (A.\text{shape} = \text{EXPANDED}) \wedge (A.\text{state} \in \{\text{FOLLOWER}, \text{ROOT}\}) \wedge (\forall B \in N(A) : B.\text{state} \neq \text{IDLE})$
  - 16:    $\wedge (A \text{ has a tail-child } B : B.\text{shape} = \text{CONTRACTED}) \rightarrow$
  - 17:   **if**  $\text{READ}(\perp, \text{state}) = \text{ROOT}$  **then**  $\text{WRITE}(\perp, \text{dir}, \text{GETNEXTDIR(counter-clockwise)}).$
  - 18:   Find a port  $p \in \text{TAILCHILDREN}()$  s.t.  $\text{READ}(p, \text{shape}) = \text{CONTRACTED}.$  ▷ See Algorithm 2.
  - 19:   Let  $p'$  be the label of the tail-child's port that will be connected to  $p$  after the pull handover.
  - 20:    $\text{WRITE}(p, \text{parent}, p').$
  - 21:    $\text{PULL}(p).$
  - 22:  $\alpha_6 : (A.\text{shape} = \text{EXPANDED}) \wedge (A.\text{state} \in \{\text{FOLLOWER}, \text{ROOT}\}) \wedge (\forall B \in N(A) : B.\text{state} \neq \text{IDLE})$
  - 23:    $\wedge (A \text{ has no tail-children}) \rightarrow$
  - 24:   **if**  $\text{READ}(\perp, \text{state}) = \text{ROOT}$  **then**  $\text{WRITE}(\perp, \text{dir}, \text{GETNEXTDIR(counter-clockwise)}).$
  - 25:    $\text{CONTRACT}(\text{TAIL}).$
-

# Asynchronous Hexagon Formation

**Theorem.** Hexagon-Formation (HF) is correct under an unfair asynchronous adversary.

Outline of analysis:

- HF is correct under an unfair sequential adversary.
- Enabled actions of HF remain enabled despite concurrent executions.
- Enabled actions of HF are executed identically in sequential and asynchronous settings.
- Any asynchronous execution of HF can be serialized.
- Any asynchronous execution of HF terminates.



# Asynchronous Hexagon Formation

The combination of:

- Correctness under an **unfair sequential adversary**,
- Enabled actions **remaining enabled** despite concurrency, and
- Enabled actions **executing identically** in sequential and asynchronous settings

immediately yields serializability and asynchronous termination, which in turn yield **asynchronous correctness**.



# A General Framework for Concurrency Control

---

Another approach to concurrency control: use **locks** to mitigate changes to an amoebot's neighborhood while it is active.

We developed a novel algorithm for **mutual exclusion** (locking) in **asynchronous, anonymous, dynamic** (moving), **constant-size memory** message passing systems.

Key Idea:

- On activation, an amoebot  $A$  first attempts to **lock its neighborhood**.
- If successful, its locked neighbors cannot move or change their memory contents.
- So  $A$  can evaluate its guards and perform its actions **as if things were sequential** (sort of).
- Failed locking attempts and expansions have no effect on the rest of the system.

Key Issue: Locks can't stop amoebots from expanding into an acting amoebot's neighborhood!

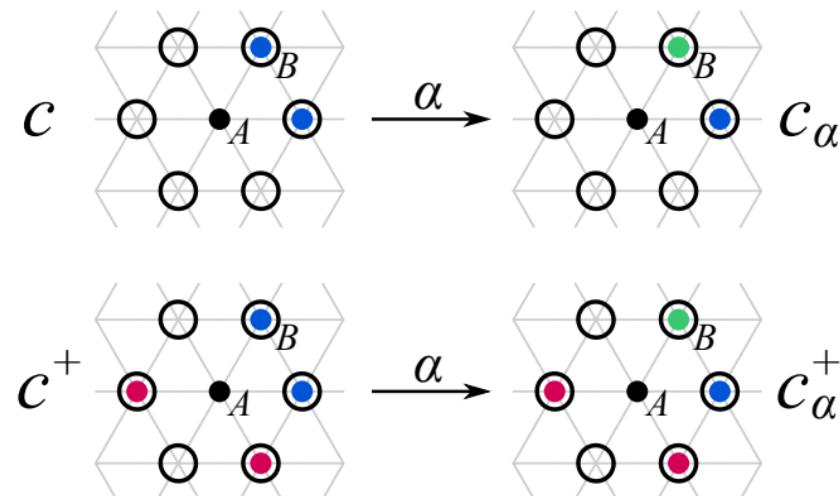
# A General Framework for Concurrency Control

We introduce a set of conventions that must be satisfied for the protocol to apply.

Convention 1: Any execution of an enabled action must succeed in the sequential setting.

Convention 2: All compute operations must precede at most one movement operation.

Convention 3: **Monotonicity**. Action executions cannot be affected by (unlocked) amoebots that concurrently enter the acting amoebot's neighborhood.



Static algorithms (i.e., those that don't involve movement) are trivially monotonic. This includes most of the leader election algorithms.

Open Question: What amoebot algorithms satisfy monotonicity?

# A General Framework for Concurrency Control

1. Validity. Any execution of an enabled action succeed in the sequential setting.
2. Computing Before Moving. Compute operations precede movement operations.
3. Monotonicity. Action executions are not affected by (unlocked) amoebots that concurrently enter the acting amoebot's neighborhood.

**Theorem.** Consider any algorithm  $\mathcal{A}$  satisfying Conventions 1-3 and let  $\mathcal{A}'$  be the algorithm obtained by the concurrency control protocol. If  $\mathcal{A}$  terminates under any **sequential execution**, then every asynchronous execution of  $\mathcal{A}'$  terminates in some sequential outcome of  $\mathcal{A}$ .

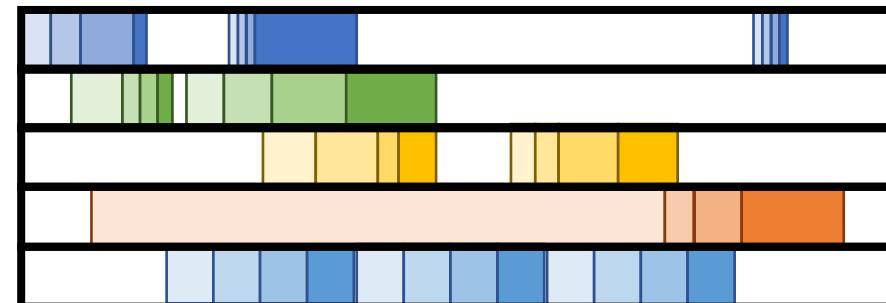
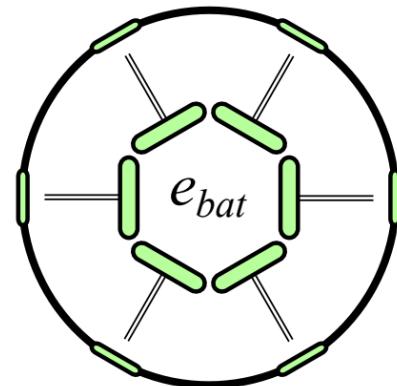


## Part II

### The Amoebot Model and its Enhancements

How can **existing algorithms** be enhanced to capture more **realistic assumptions**?

By satisfying certain conventions, **energy-agnostic, sequential** algorithms  
can be made **energy-constrained and asynchronous**.

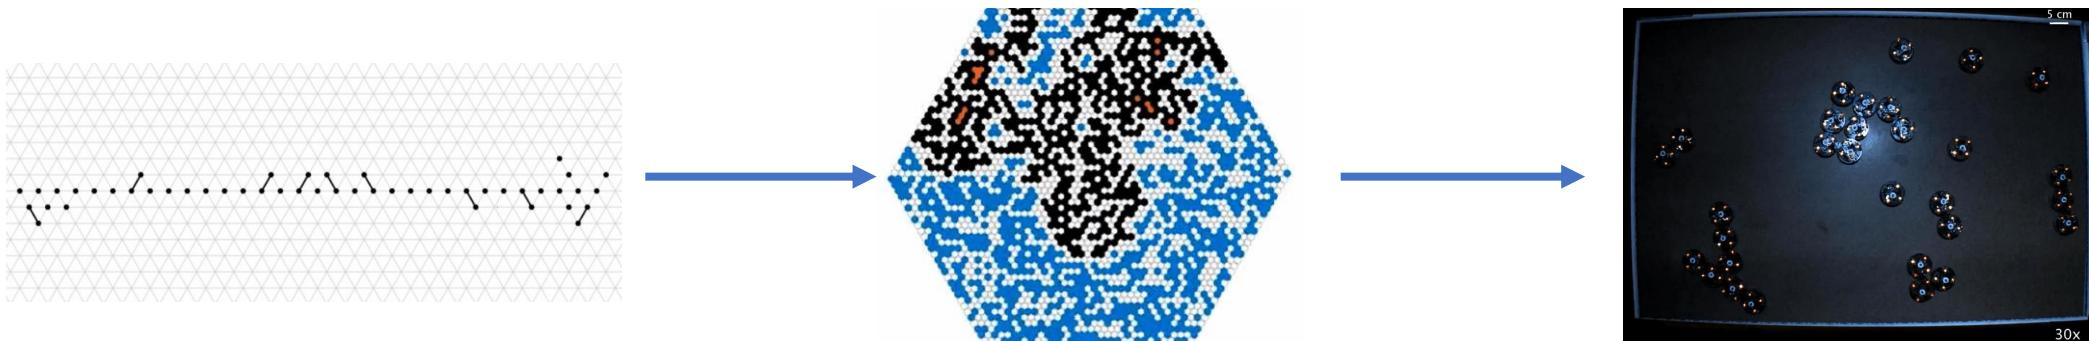


# Part III

## Stochastic Distributed Algorithms & Their Applications to Swarm Robotics and Granular Active Matter

How can **digital algorithms** be translated for simple, **analog (passive) systems**?

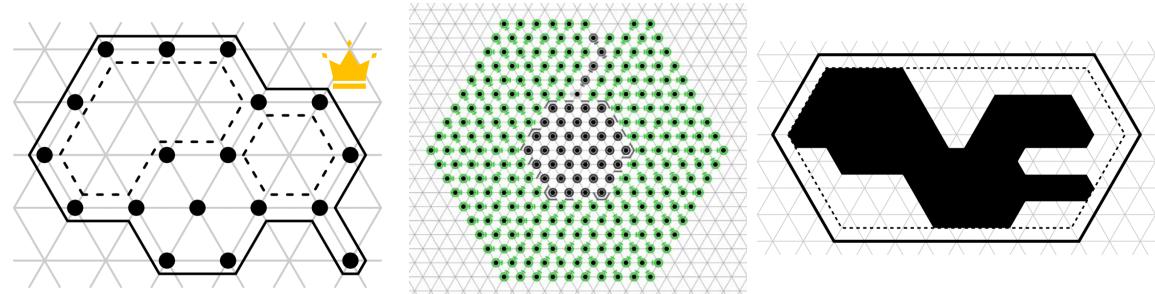
Connect **biased random decisions** to the **physics of local interactions**.



# Conclusion: Algorithmic Foundations of Programmable Matter

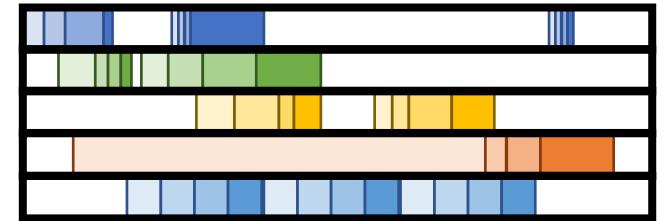
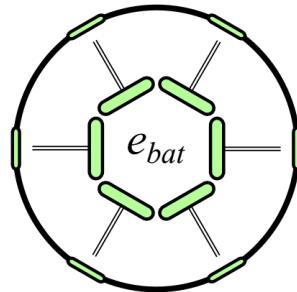
1. What are the **minimum individual capabilities** necessary to achieve **system behavior X**?

Constant-size memory, communication, and local movements suffice for complex behaviors.



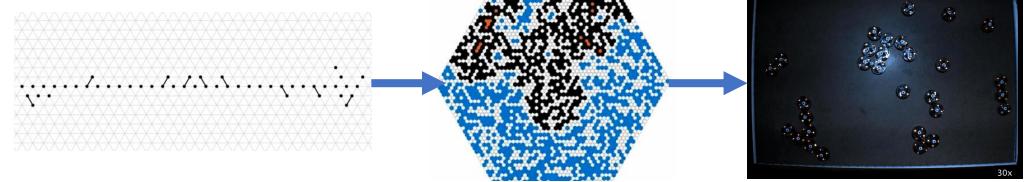
2. How can **existing algorithms** be enhanced to capture more **realistic assumptions**?

By satisfying certain conventions, energy-agnostic, sequential algorithms can be made energy-constrained and asynchronous.

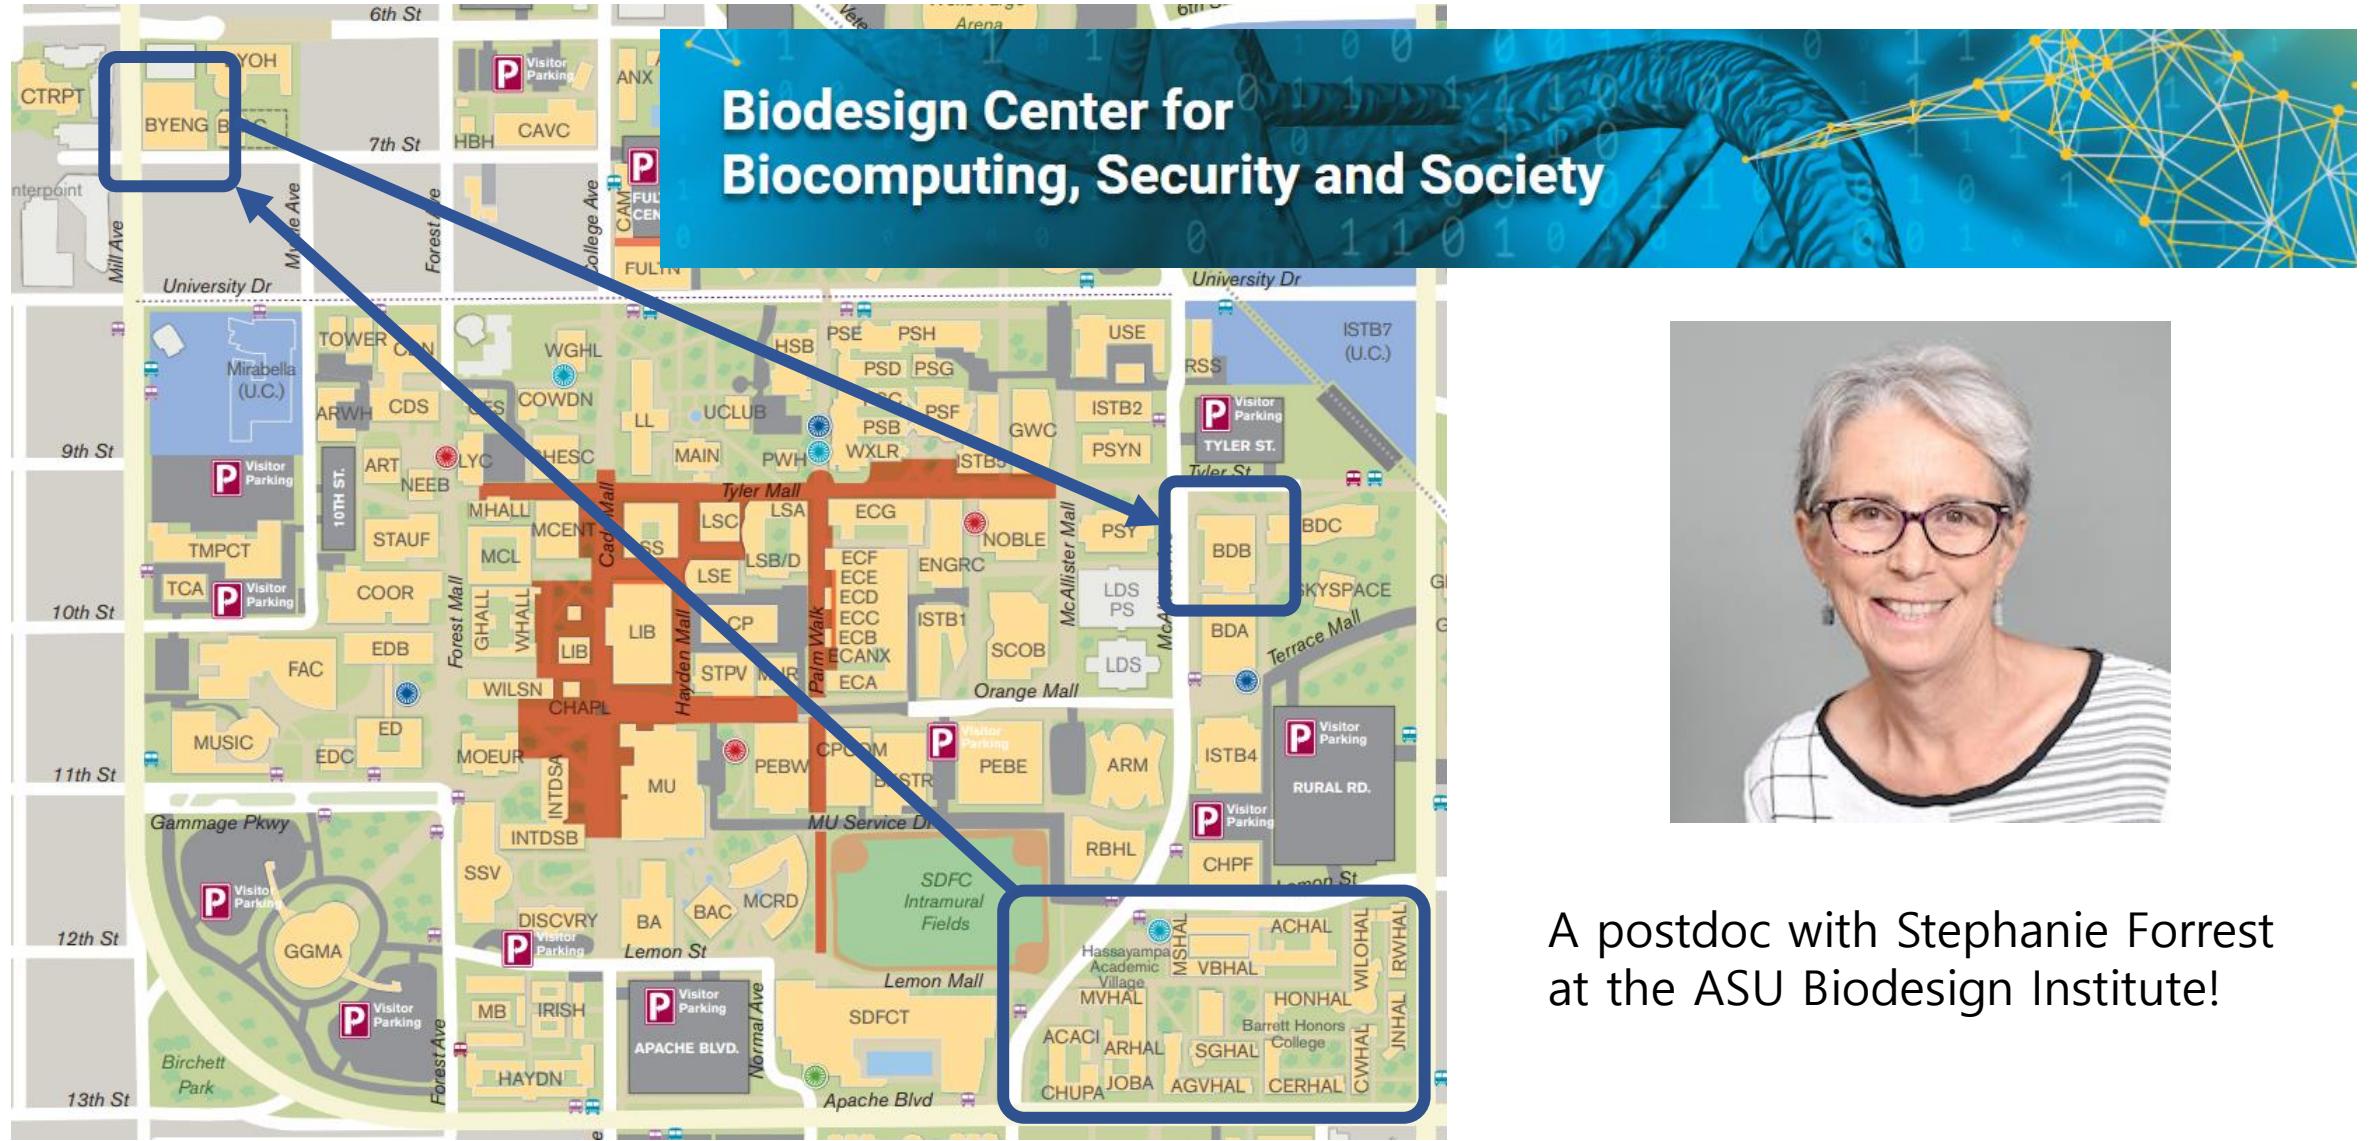


3. How can **digital algorithms** be translated for simple, **analog (passive) systems**?

Connect biased random decisions to the physics of local interactions.



# What's Next?



# I've Got Some People Who Carry Me

---



# Thank you!

[sops.engineering.asu.edu](http://sops.engineering.asu.edu)

[jdaymude.github.io](http://jdaymude.github.io)



# List of Publications: Dissertation (Chronologically)

---

Paper	Conference	Journal
"A Markov Chain Algorithm for Compression in Self-Organizing Particle Systems." Cannon, D., Randall, Richa.	PODC 2016	In Preparation
"A Stochastic Approach to Shortcut Bridging in Programmable Matter." Andrés Arroyo, Cannon, D., Randall, Richa.	DNA 2017	Natural Computing
"On the Runtime of Universal Coating for Programmable Matter." D., Derakhshandeh, Gmyr, Porter, Richa, Scheideler, Strothmann.		Natural Computing
"Improved Leader Election for Self-Organizing Programmable Matter." D., Gmyr, Richa, Scheideler, Strothmann.	ALGOSENSORS 2017	
"Phototactic Supersmarticles." Savoie, Cannon, D., Warkentin, Li, Richa, Randall, Goldman.		Artificial Life and Robotics
"A Local Stochastic Algorithm for Separation in Heterogeneous Self-Organizing Particle Systems." Cannon, D., Gökm̄en, Randall, Richa.	PODC 2018 (BA) RANDOM 2019	In Preparation
"Convex Hull Formation for Programmable Matter." D., Gmyr, Hinnenthal, Kostitsyna, Scheideler, Richa.	ICDCN 2020	
"Bio-Inspired Energy Distribution for Programmable Matter." D., Richa, Weber.	ICDCN 2021	

# List of Publications: Dissertation (Chronologically)

---

Paper	Conference	Journal
"Programming Active Granular Matter with Mechanically Included Phase Changes." Li, Dutta, Cannon, D., Avinery, Aydin, Richa, Goldman, Randall.		Science Advances
"The Canonical Amoebot Model: Algorithms and Concurrency Control." D, Richa, Scheideler.*	In Preparation	

# List of Publications: Non-Dissertation (Chronologically)

---

Paper	Conference	Journal
"Computing by Programmable Particles." D., Hinnenthal, Richa, Scheideler. Book Chapter in <u>Distributed Computing by Mobile Entities</u> .		2018
"Simulation of Programmable Matter Systems Using Active Tile-Based Self-Assembly." Alumbaugh, D., Demaine, Patitz, Richa.	DNA 2019	Natural Computing*
"Preventing Extreme Polarization of Political Attitudes." Axelrod, D., Forrest.		PNAS*
"Mutual Exclusion for Asynchronous, Anonymous, Dynamic, Constant-Size Memory Message Passing Systems." D., Scheideler, Richa.	DISC 2021**	
"AmoebotSim: A Visual Simulator for the Amoebot Model of Programmable Matter." D., Gmyr, Hinnenthal.**		
"Aggregation Without Computation: Negative Results and a Noisy, Discrete Adaptation." D., Harasha, Richa, Yiu.**		

\*Under review.

\*\*Manuscript in preparation.